

Deep Learning for beginners

Part I - Fundamentals

JB Fiche, CBS-Montpellier & Plateforme MARS-MRI

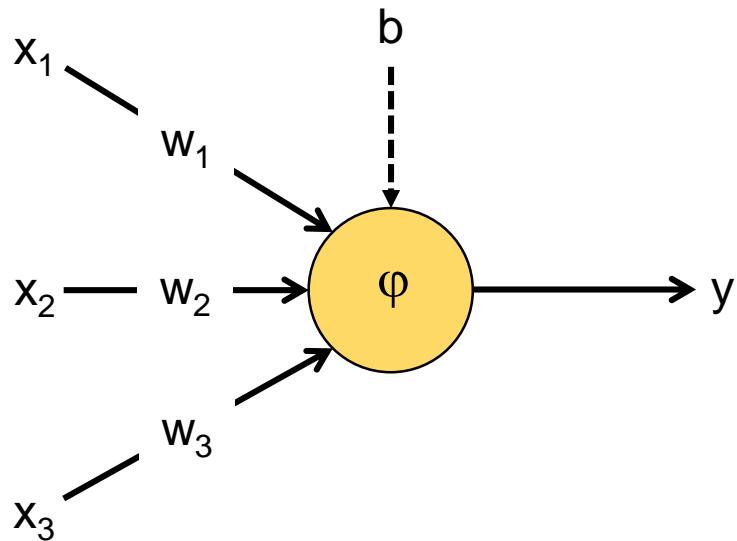
Francesco Pedaci, CBS-Montpellier

Volker Bäcker, CRBM & MRI

Cédric Hassen-Khodja, CRBM & MRI

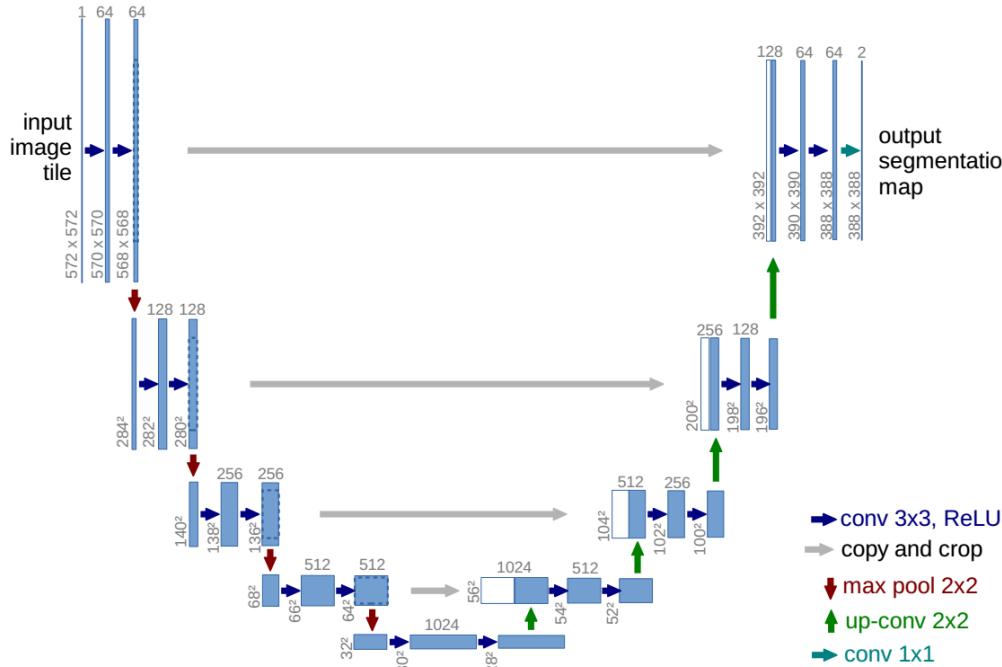
End-of-the-day goal

- Understand the differences between **machine-learning & deep learning**
- Understand what an **Artificial Neuron Network (ANN)** is and what are the main parameters to characterize them
- What are the **fundamentals for building and training an ANN**



End-of-the-day goal

- What is a **Convolutional Neural Network (CNN)** and why they are useful for image analysis
- Have seen a **few applications of Deep Learning** for image analysis and several of the most common architectures.

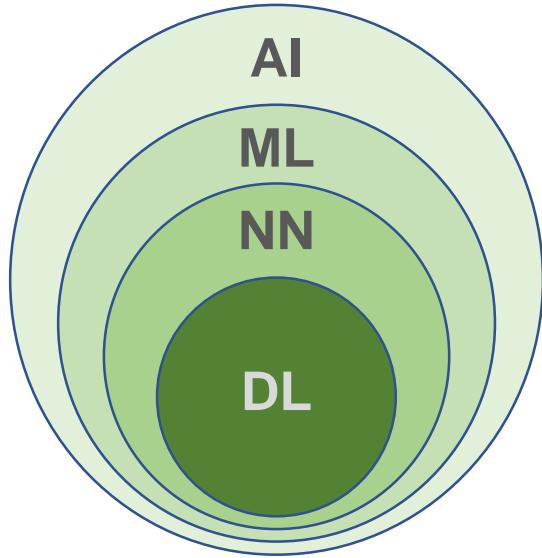


Outline

I. Introduction

- i. Difference between Machine Learning (ML) & Deep Learning (DL)
- ii. Brief historical overview
- iii. Applications for image analysis
- iv. When using DL and which tools are available?

Machine Learning vs. Deep Learning

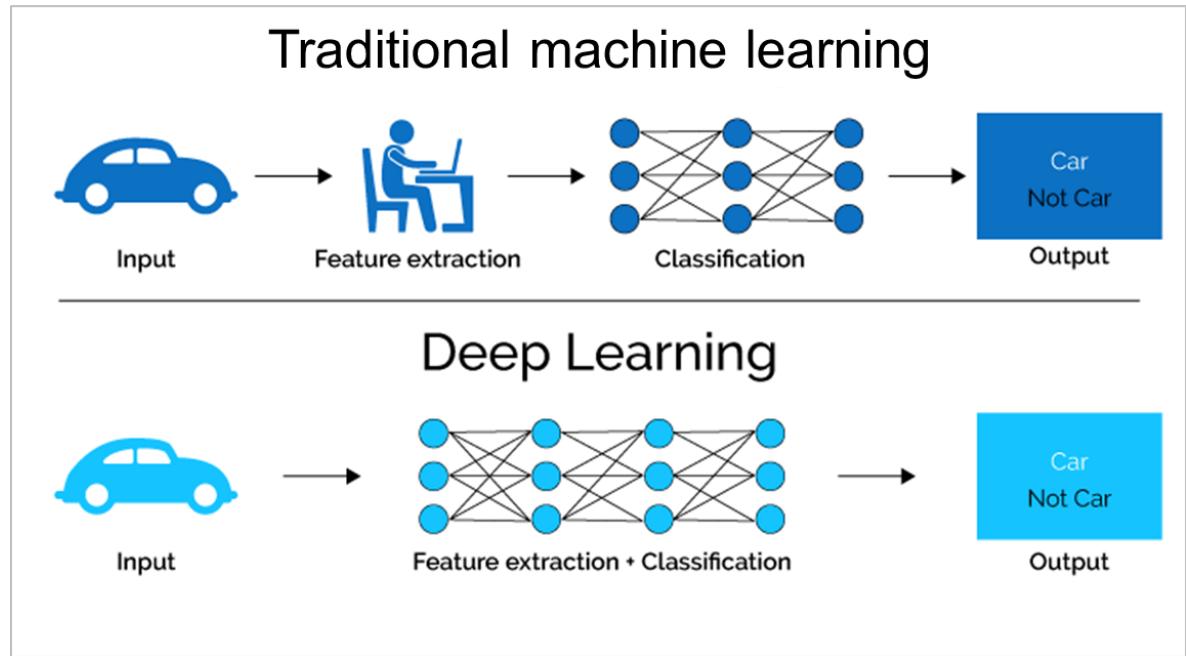


AI = artificial intelligence

ML = machine learning

NN = neural network

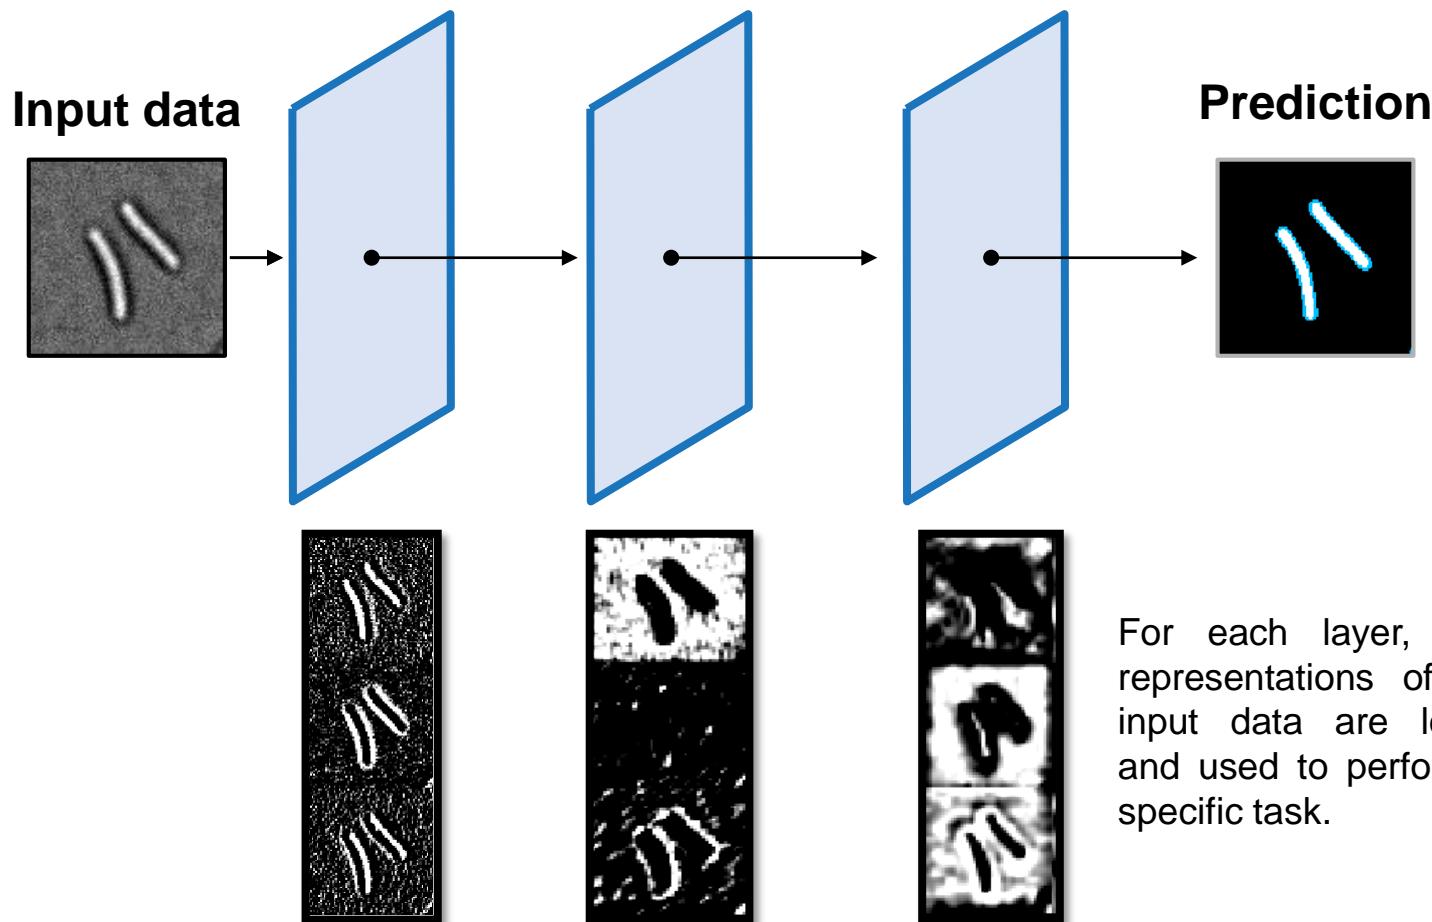
DL = deep learning



Pic Credit: Xenonstack | Machine Learning vs Deep Learning

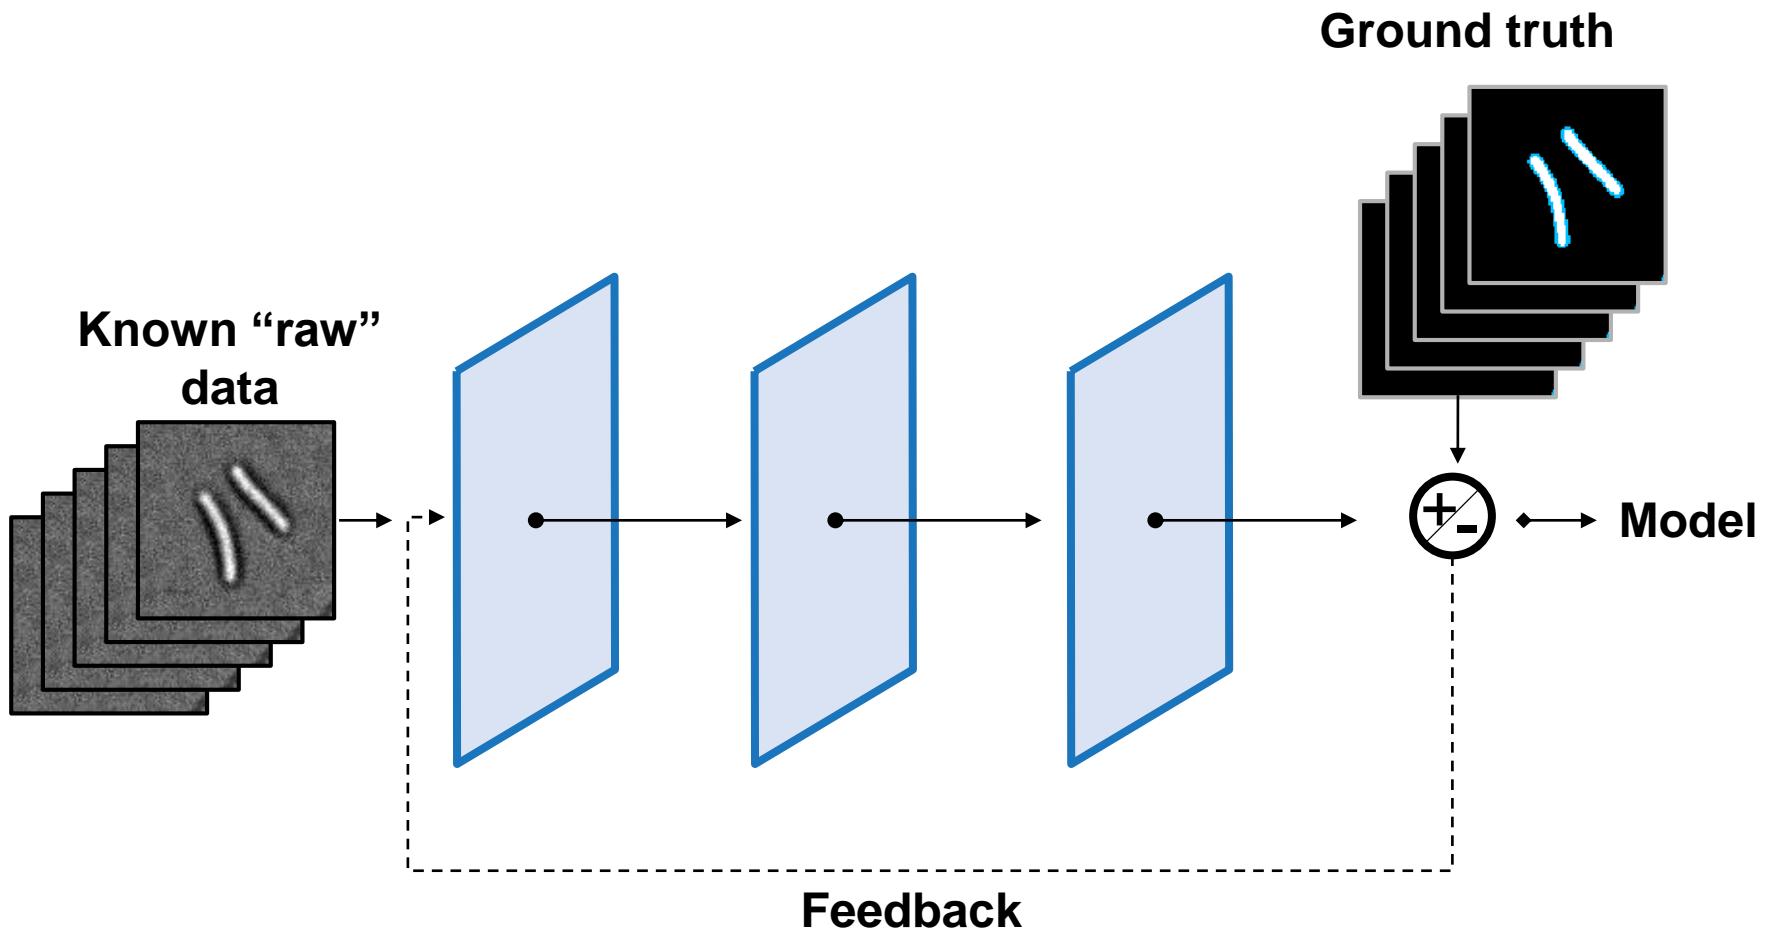
Machine Learning algorithms are searching for a useful representations of the input data.

Deep Learning : why “Deep” ?



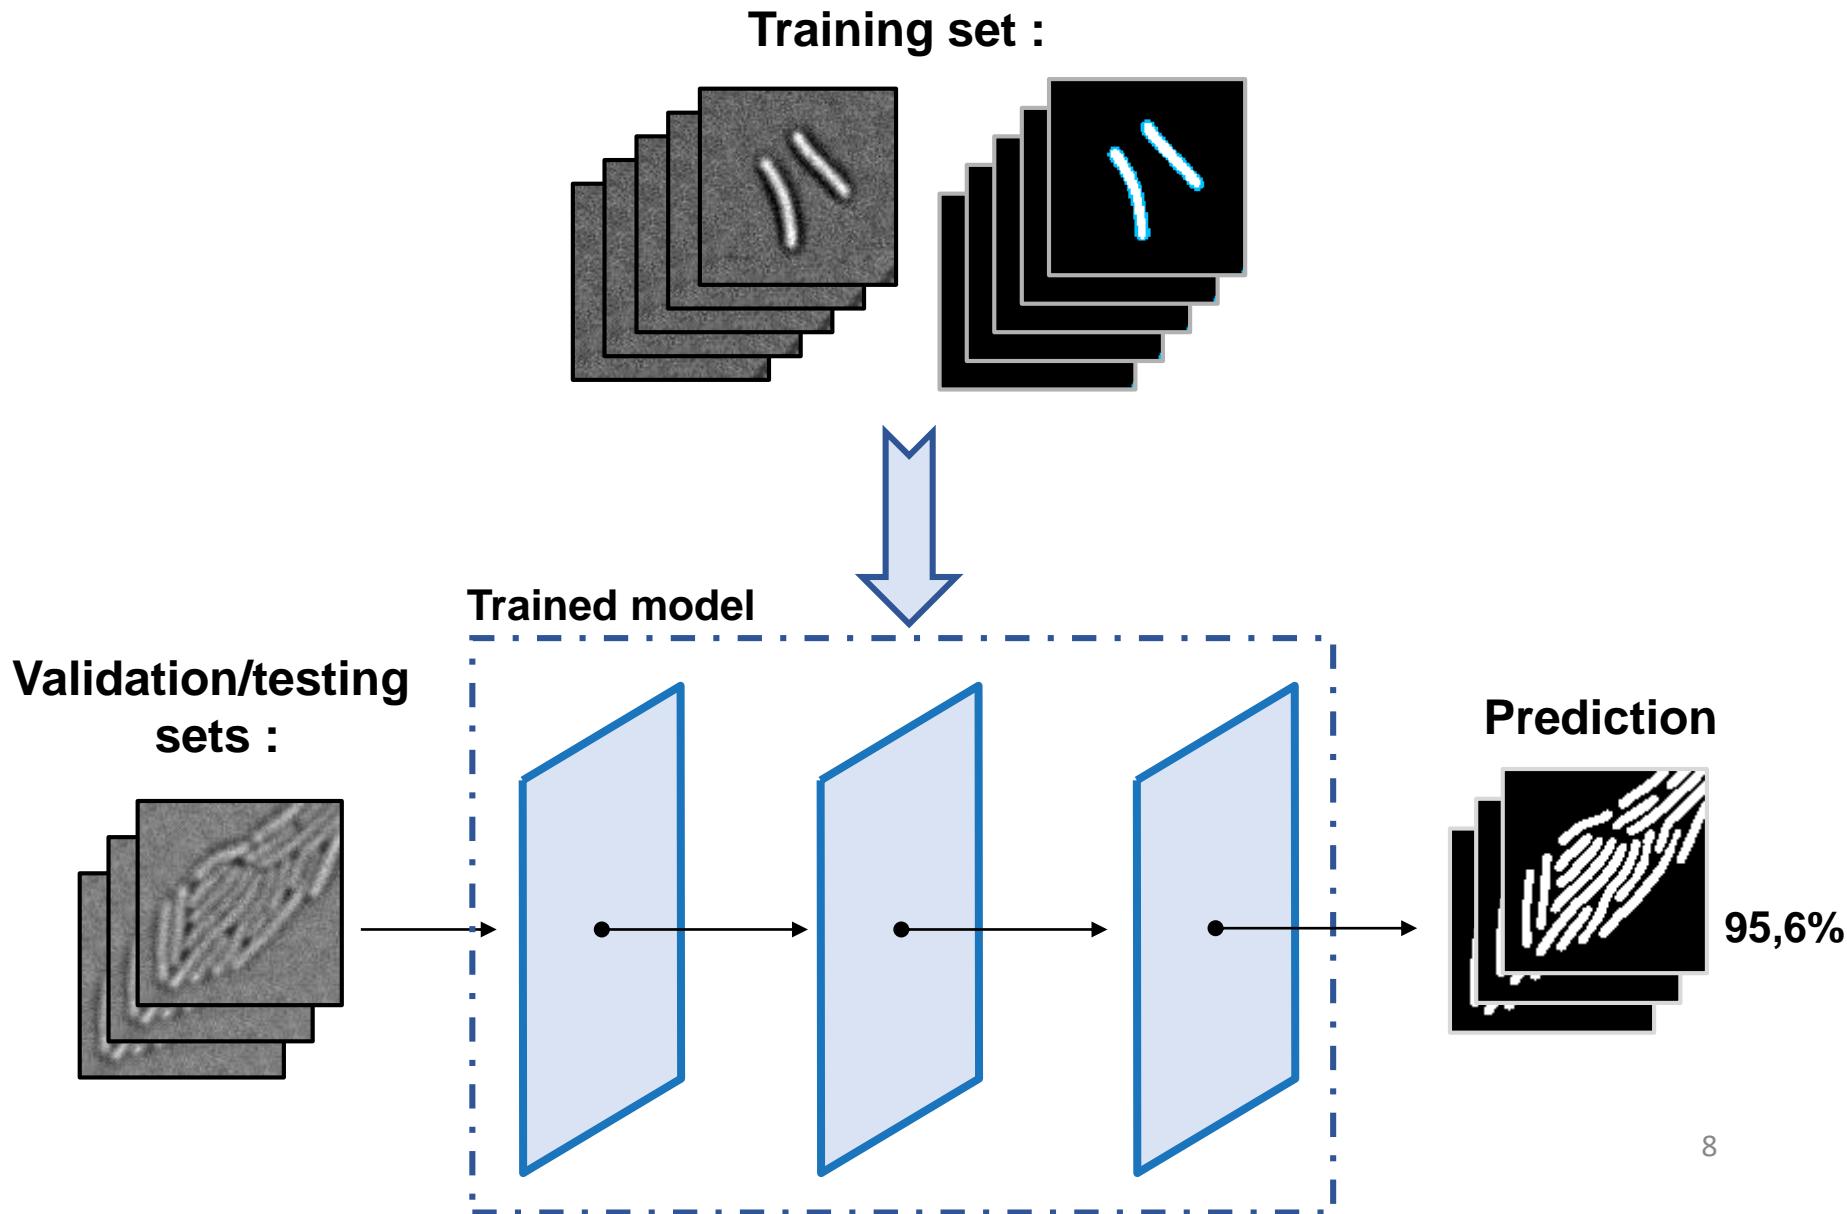
The « Deep » in Deep Learning is a reference to the **successive layers of representations** the network is using to perform its task.

Supervised Deep Learning : what “training” means?



During the training, the network will learn how to represent the data in a useful way in order to perform a specific task (segmentation, denoising, etc.)

Supervised Deep Learning : what “training” means?

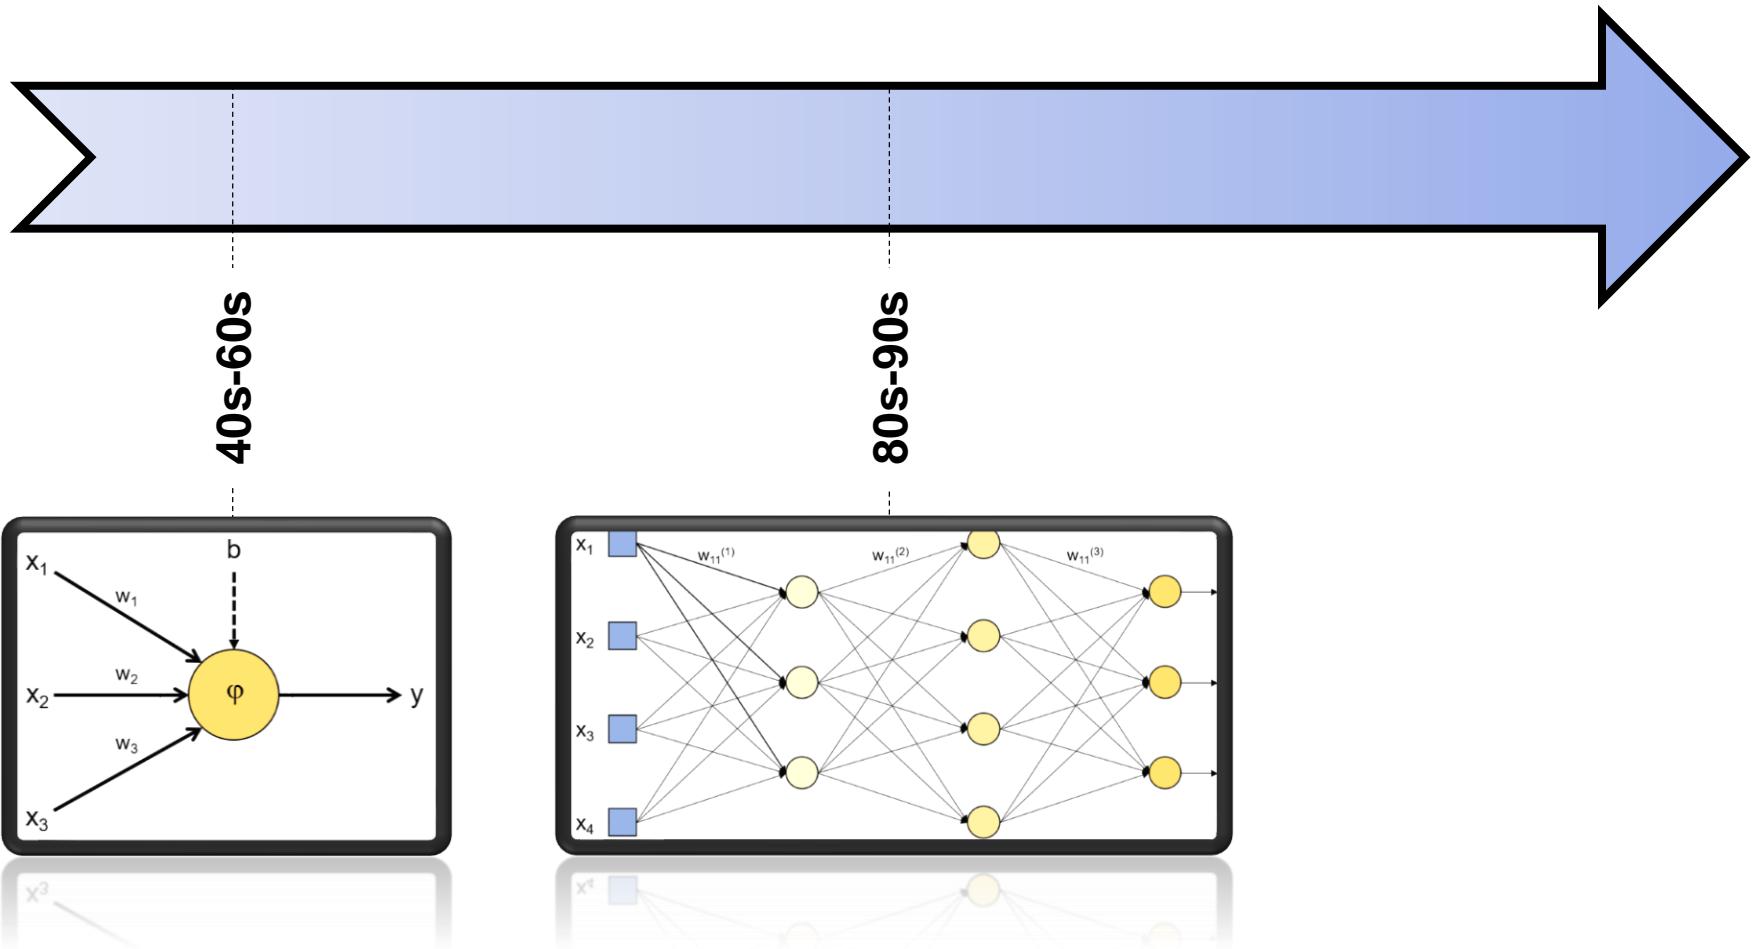


Outline

I. Introduction

- i. Difference between Machine Learning (ML) & Deep Learning (DL)
- ii. Brief historical overview**
- iii. Applications for image analysis
- iv. When using DL and which tools are available?

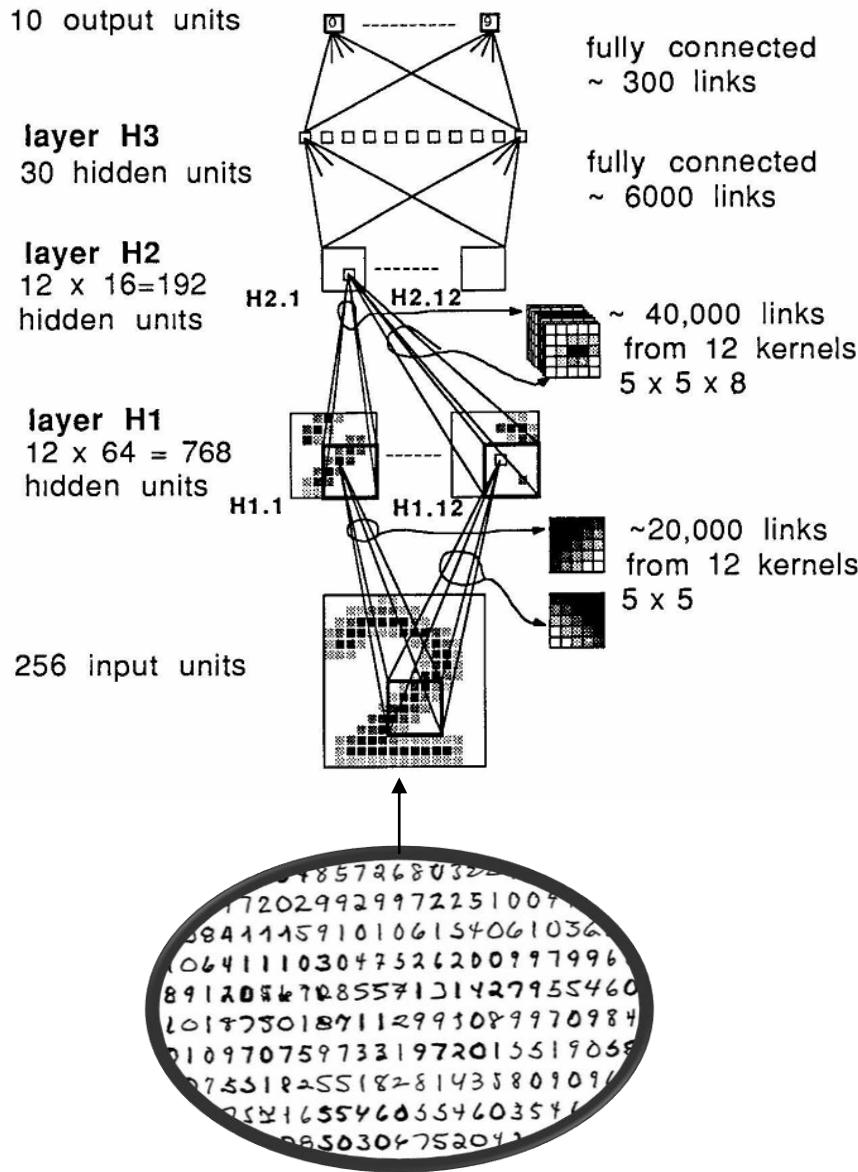
Brief historical overview



Cybernetics : linear
neural network models
able to handle simple
classification tasks

Connectivism : Development of
the **back-propagation** algorithm
for the training of deep neural
network. First **convolutional**
network.

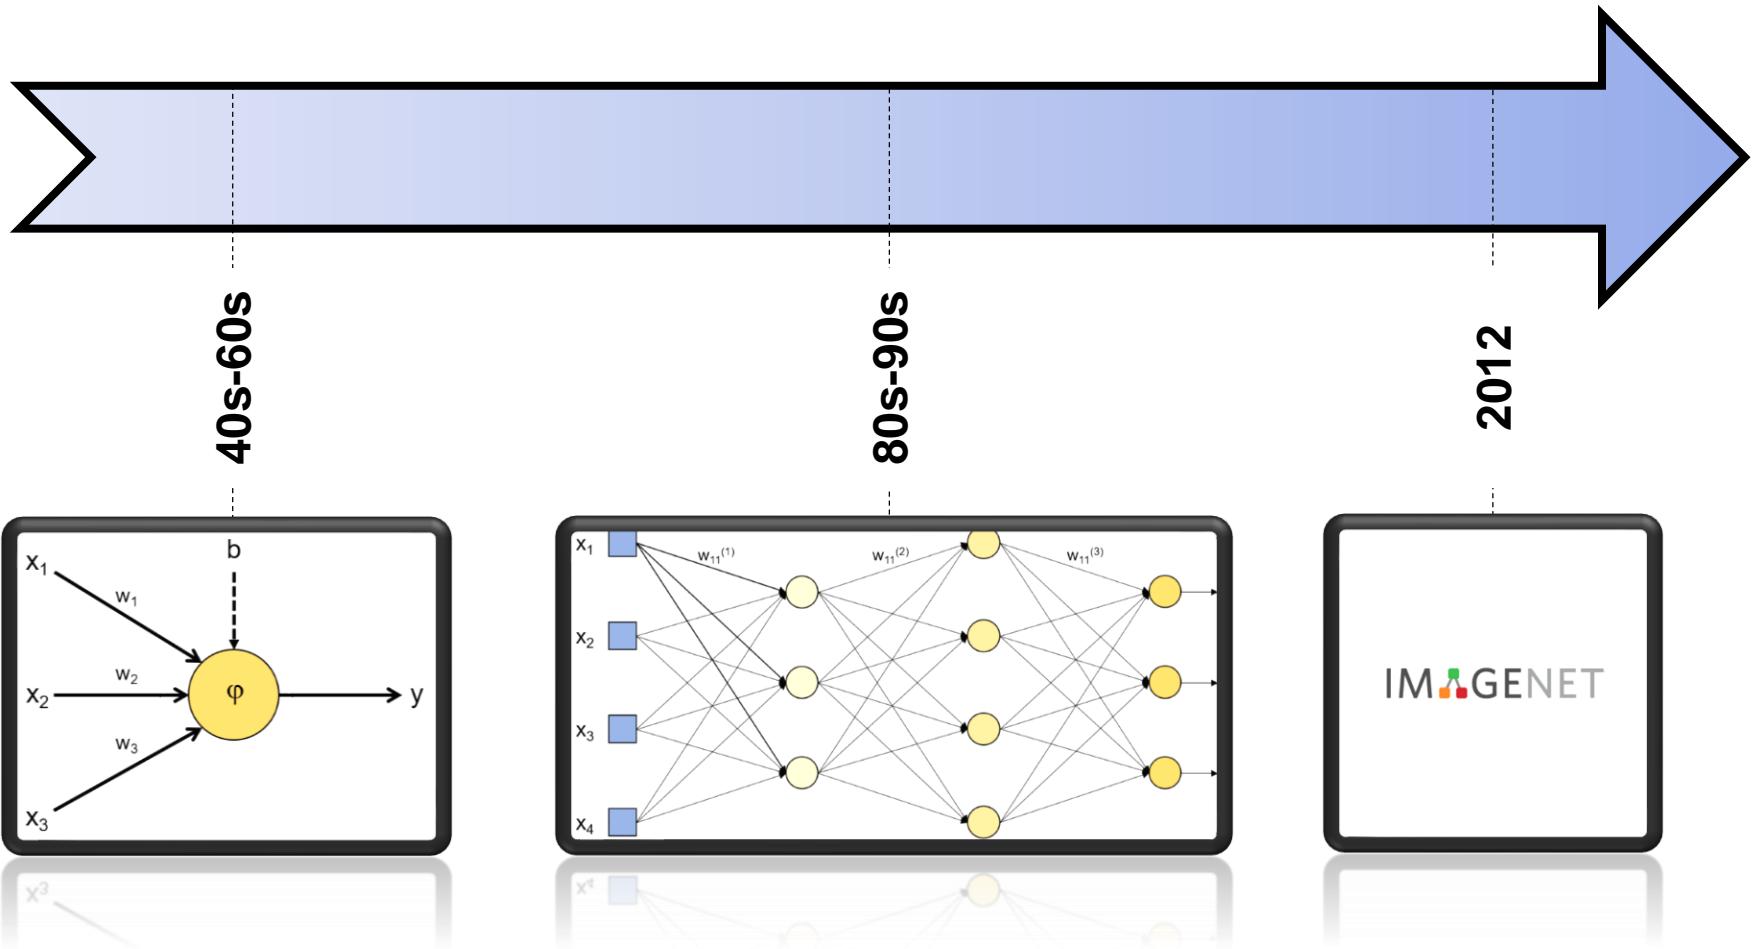
Hand-written zipcode classification - 1989



Deep learning models were already successfully used in the 90's for complicated tasks such as hand-written digit classification.

The same algorithms are still used nowadays for image analysis.

Brief historical overview

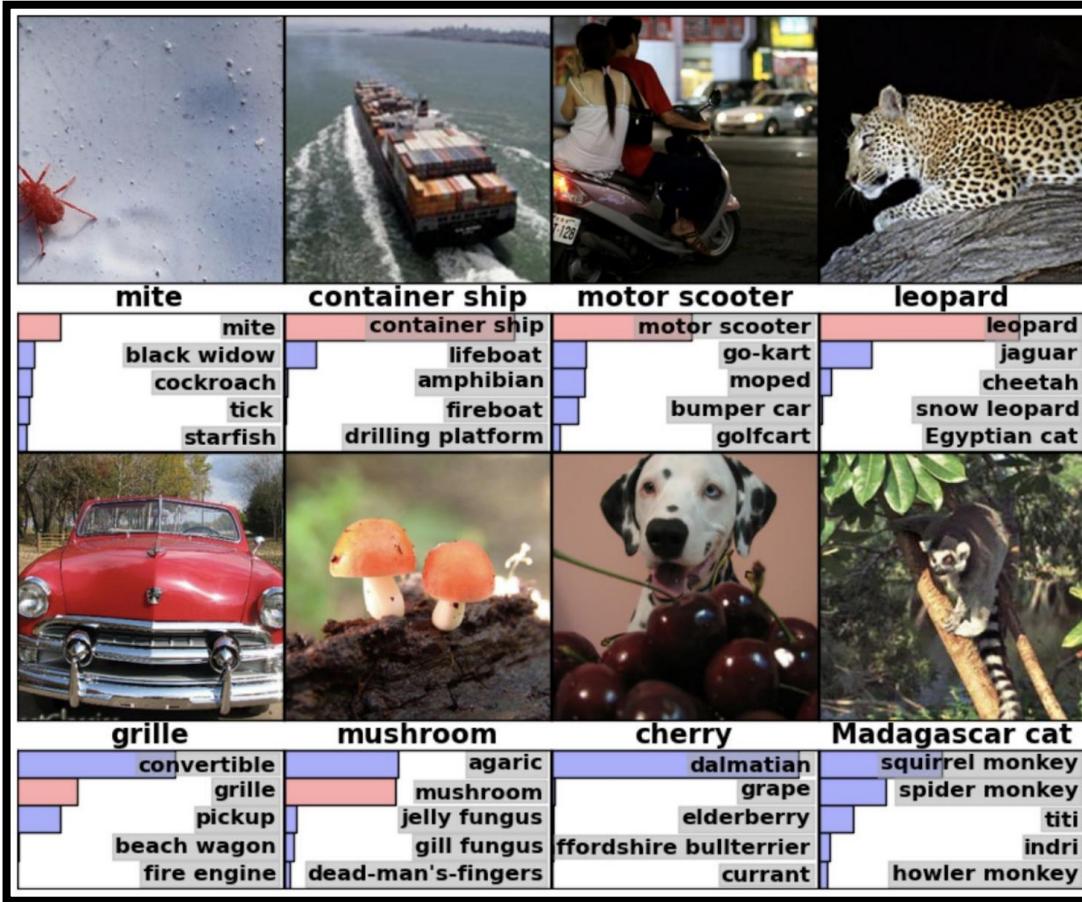


Cybernetics : linear
neural network models
able to handle simple
classification tasks

Connectivism : Development of
the **back-propagation** algorithm
for the training of deep neural
network. First **convolutional**
network.

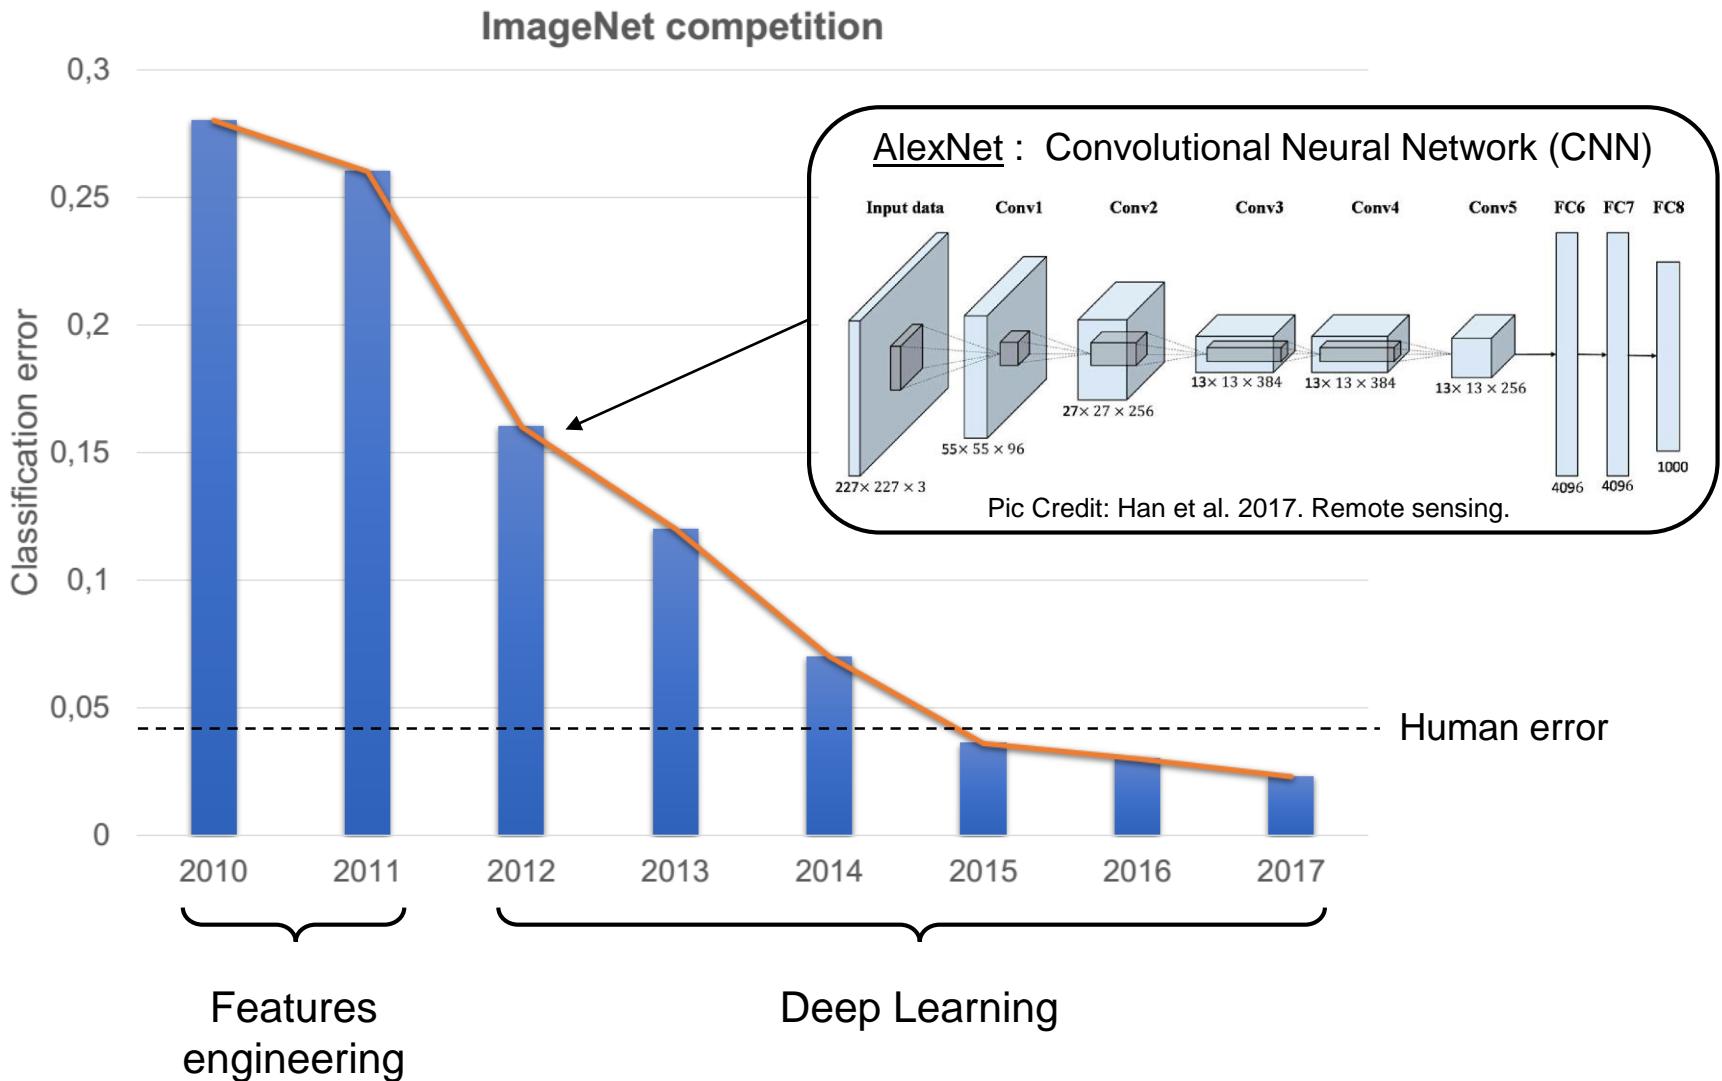
First time a **Deep**
Learning convolutional
algorithm wins the
ImageNet competition

2012 : Breakthrough for image classification

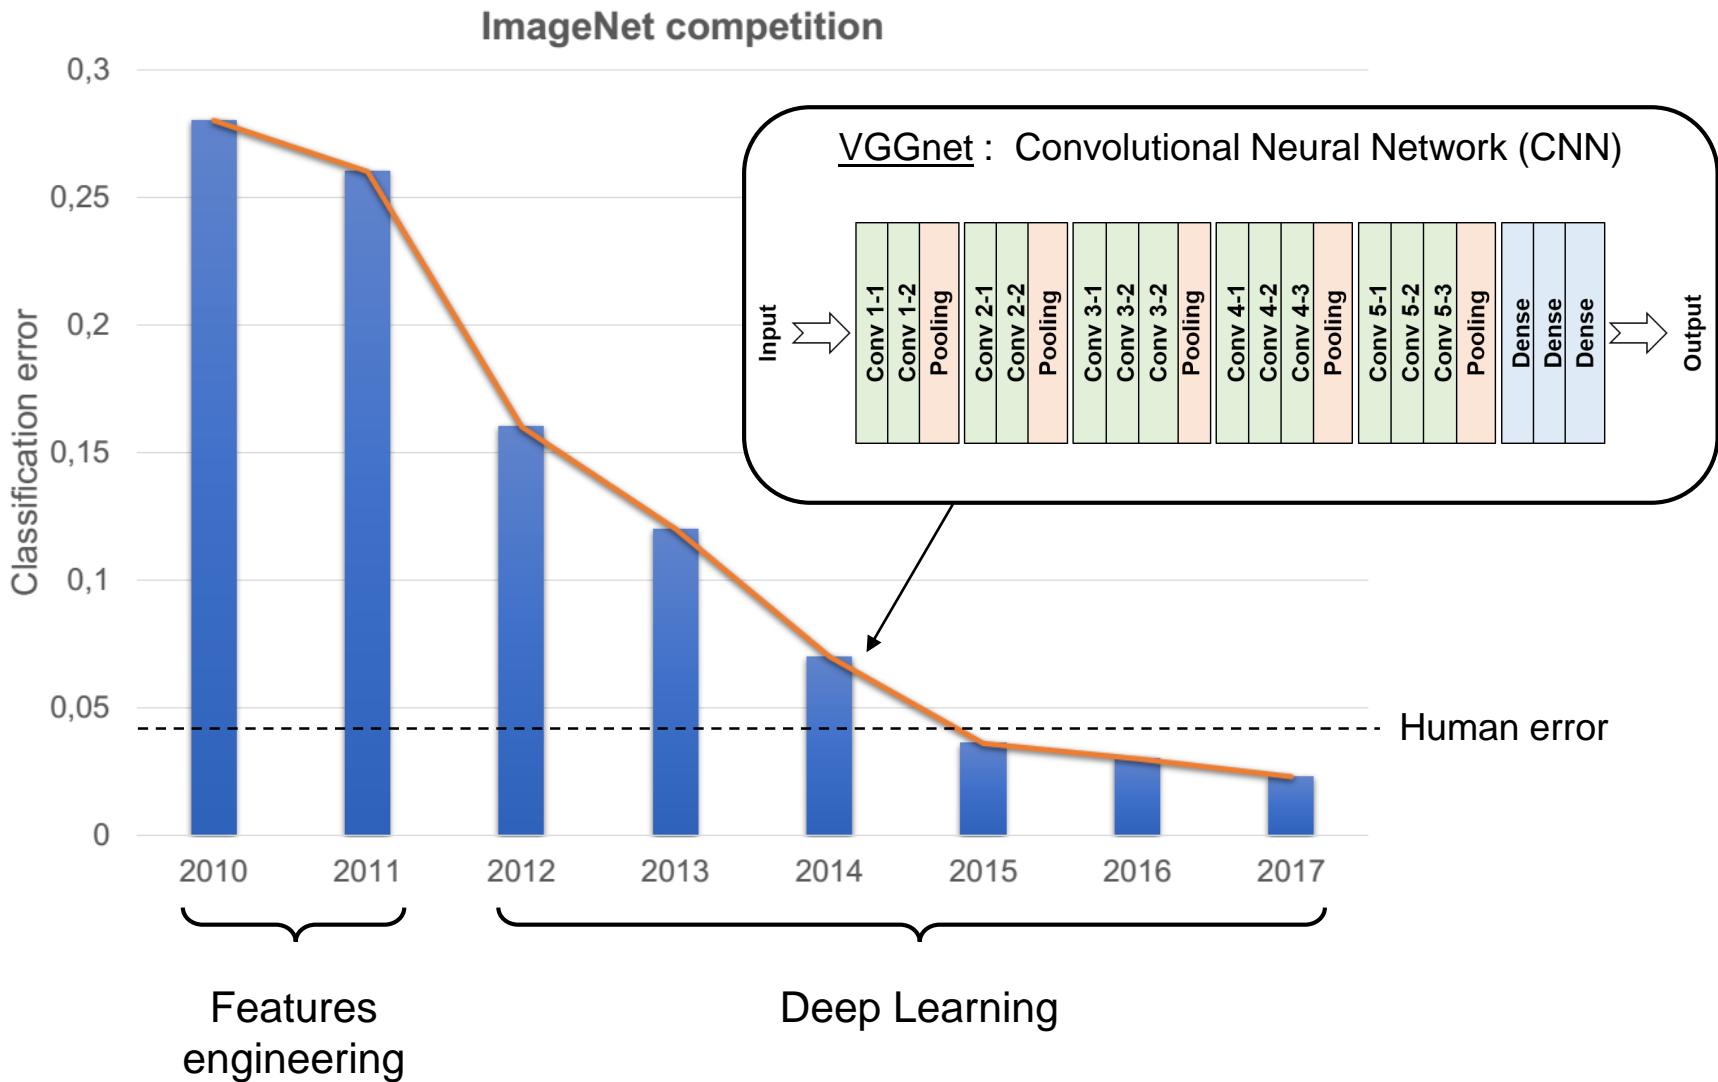


The Image Large Scale Visual Recognition Challenge started in 2010 using the ImageNet database (15M labelled images and 22,000 different object categories).

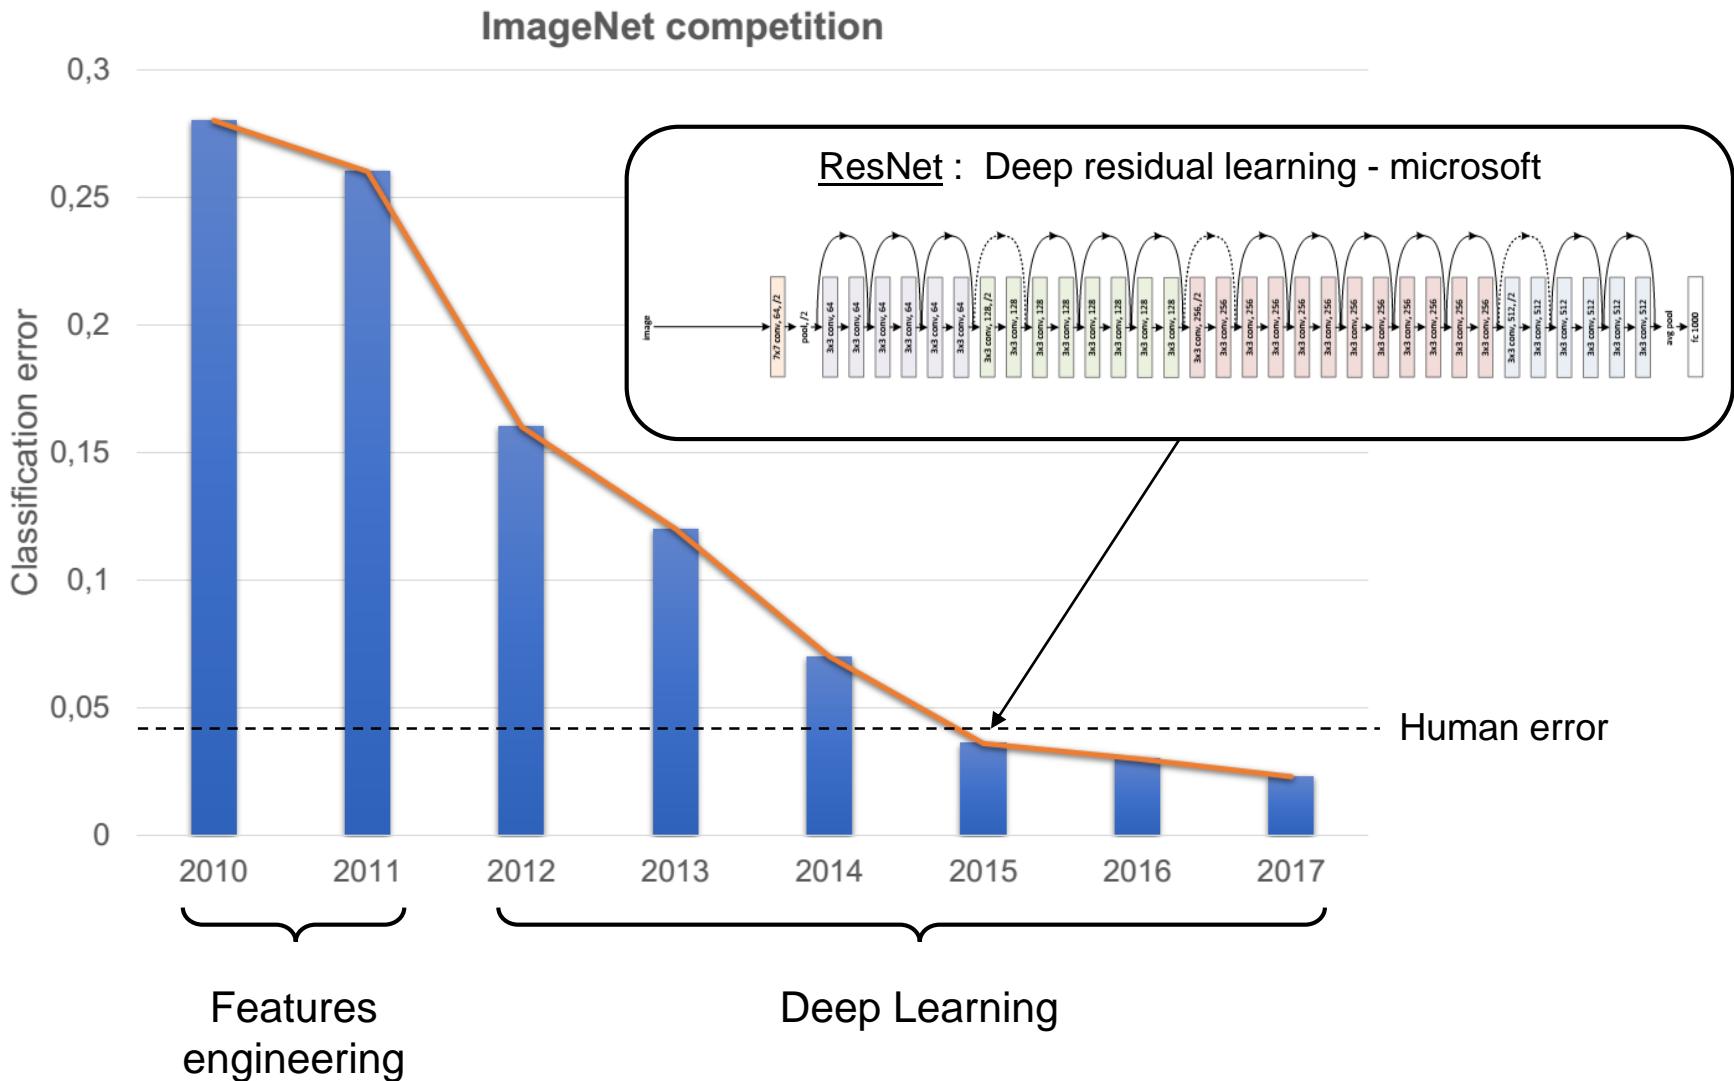
2012 : Breakthrough for image classification



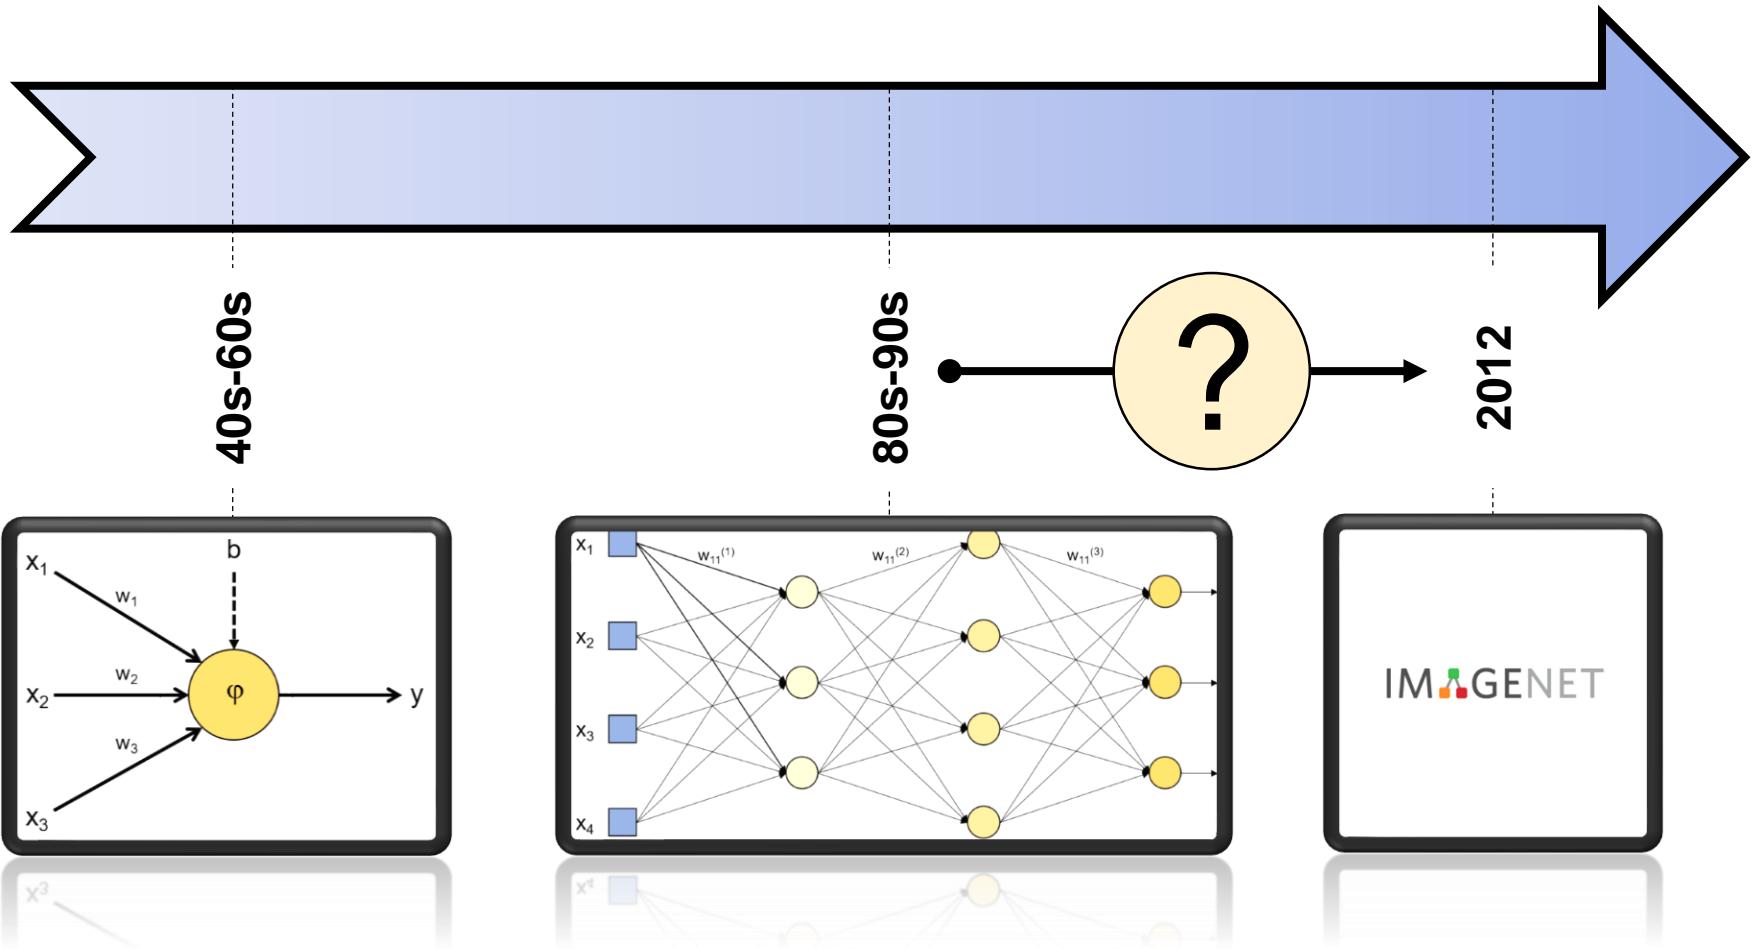
2012 : Breakthrough for image classification



2012 : Breakthrough for image classification



Why only now?

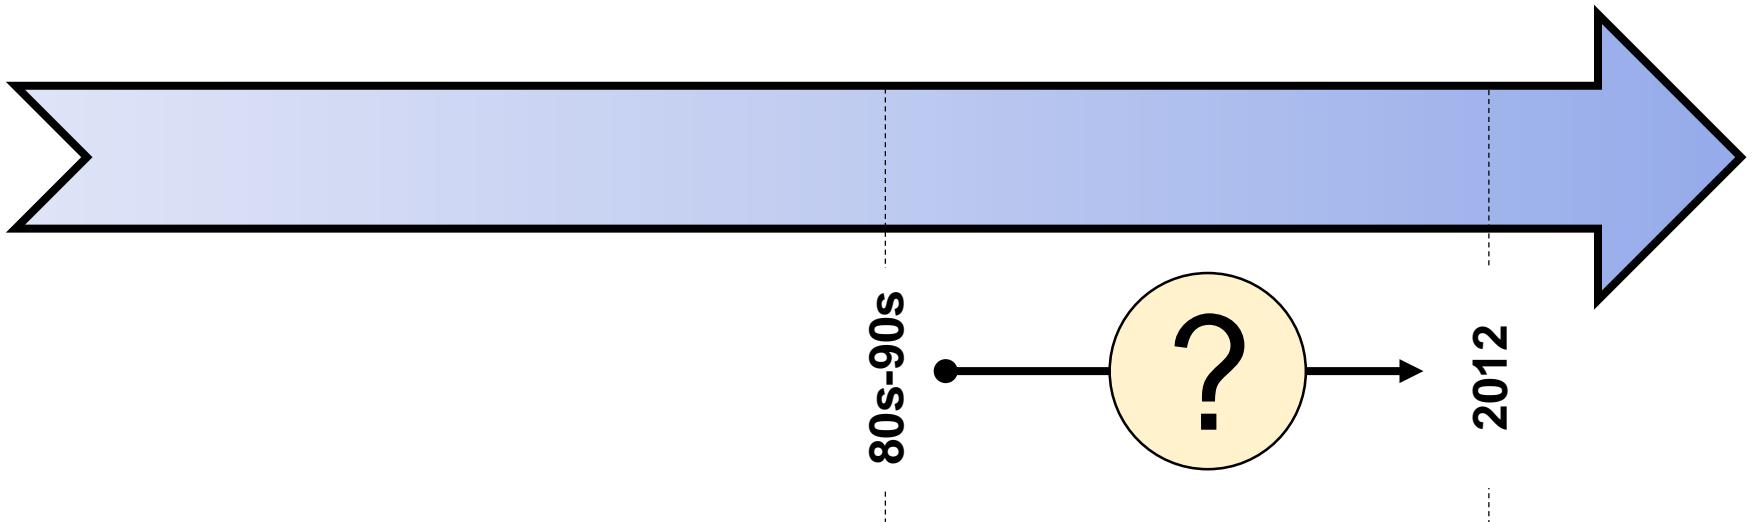


Cybernetics : linear
neural network models
able to handle simple
classification tasks

Connectivism : Development of
the **back-propagation** algorithm
for the training of deep neural
network. First **convolutional**
network.

First time a **Deep**
Learning convolutional
algorithm wins the
ImageNet competition

Why only now?



1- Internet : offers a huge source of **high-resolution annotated images**, essential for an efficient **training**.



2- Computer resources : faster CPU and GPU, able to perform $>10^{12}$ **operations/s**



NVIDIA TESLA
 $>3000\text{€}$

3- New algorithms : activation function (ReLU), faster optimizers, methods to avoid overfitting (Dropout), etc.

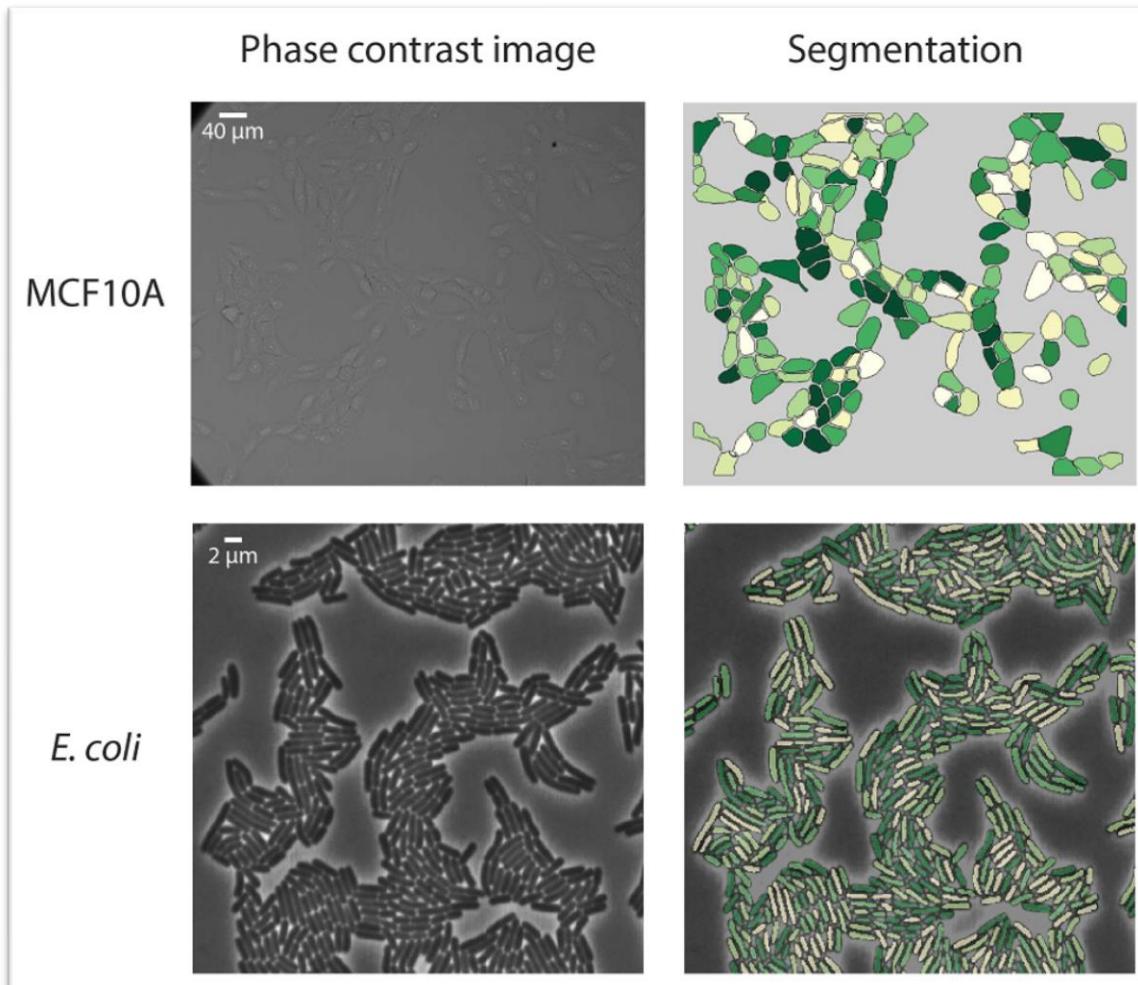
Outline

I. Introduction

- i. Difference between Machine Learning (ML) & Deep Learning (DL)
- ii. Brief historical overview
- iii. Applications for image analysis**
- iv. When using DL and which tools are available?

Deep Learning applications : a few examples

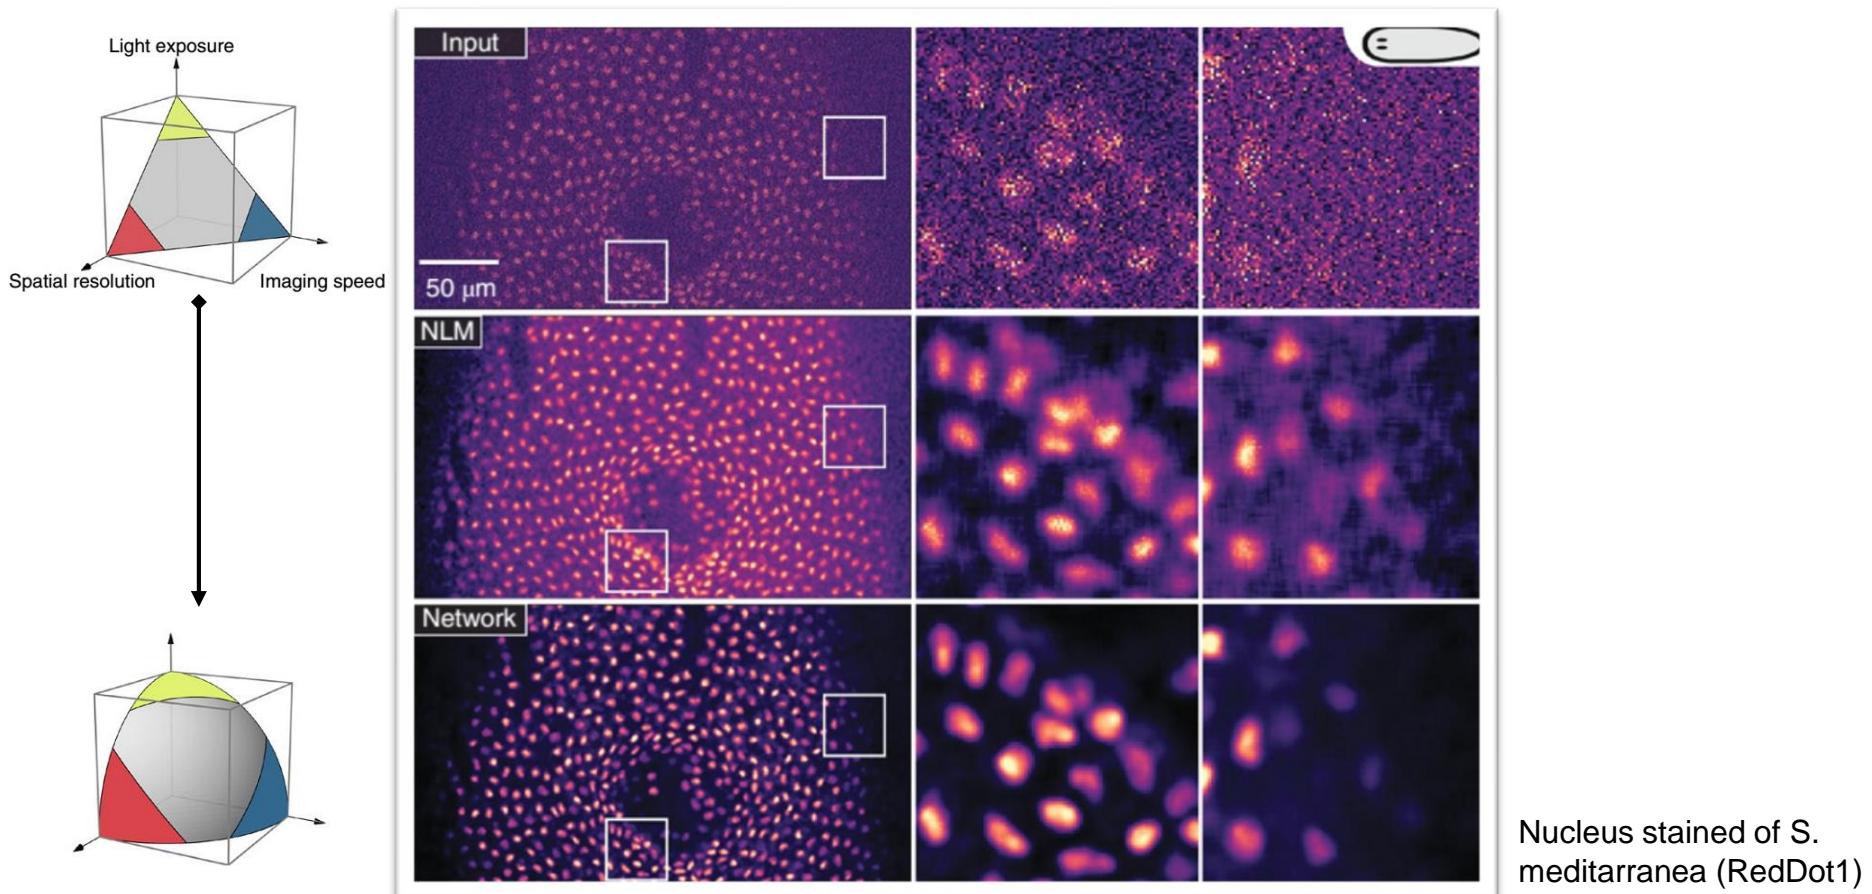
I. Segmentation



Deep Learning applications : a few examples

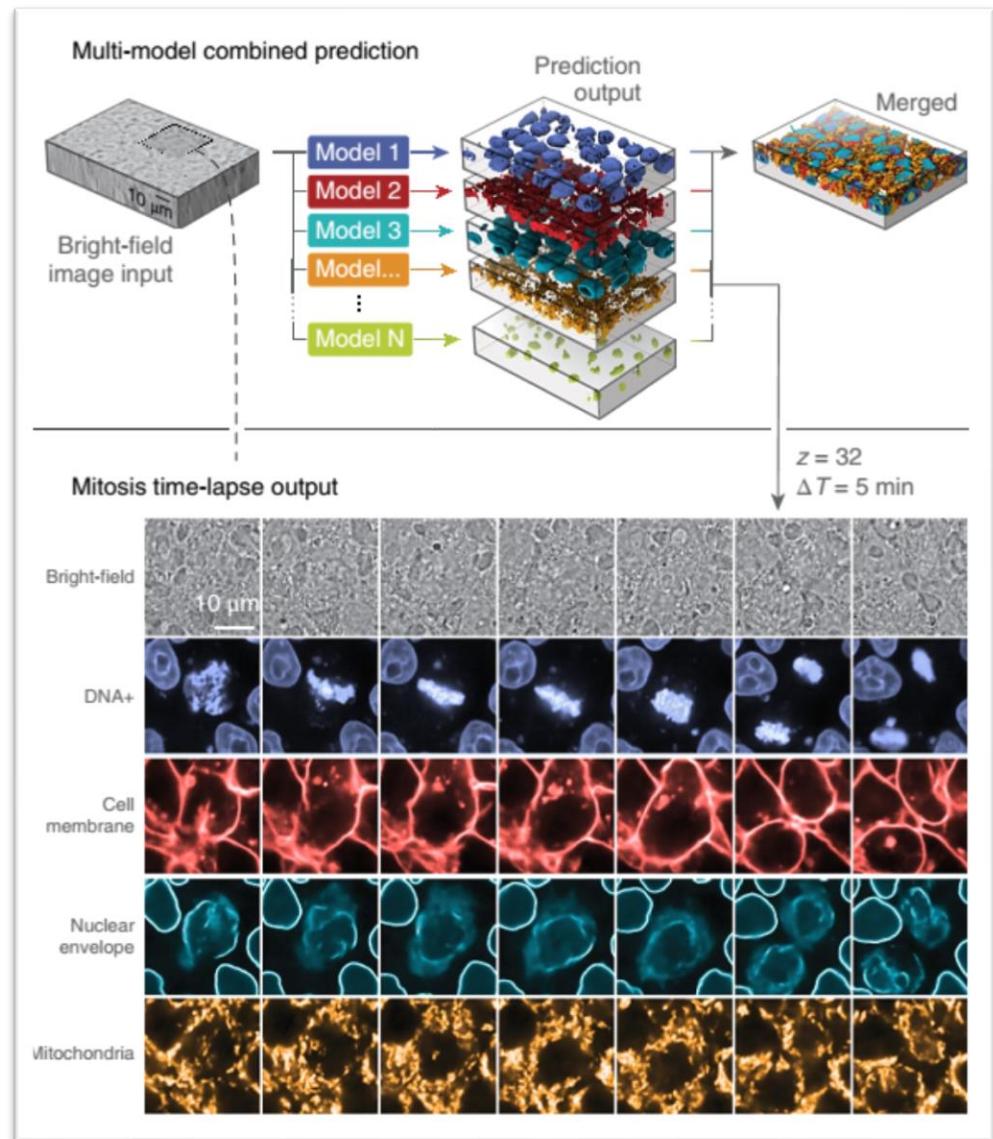
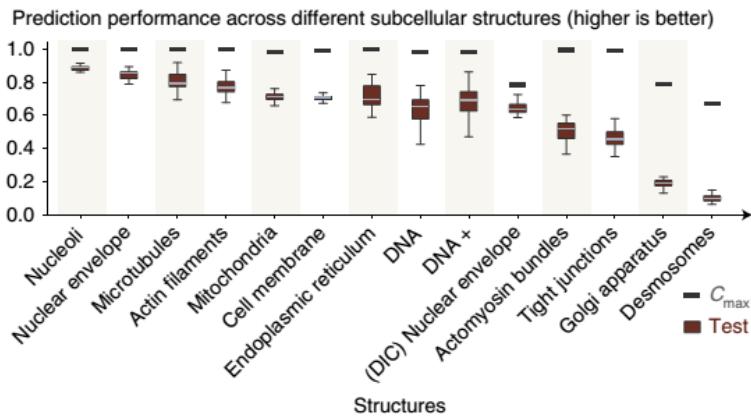
I. Segmentation

II. Image restoration



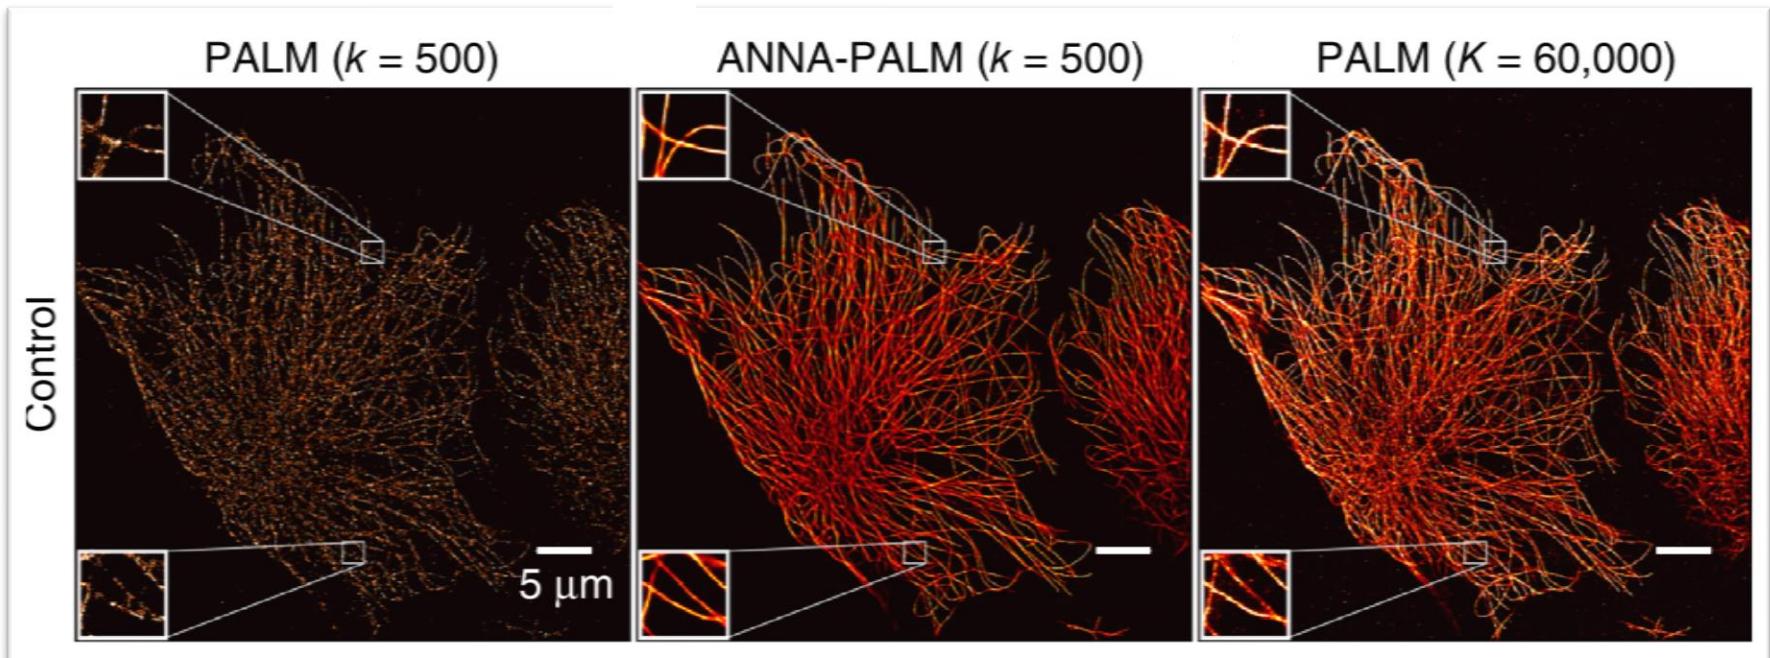
Deep Learning applications : a few examples

- I. Segmentation
- II. Image restoration
- III. Image prediction



Deep Learning applications : a few examples

- I. Segmentation
- II. Image restoration
- III. Image prediction
- IV. Super-resolution prediction**



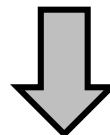
Outline

I. Introduction

- i. Difference between Machine Learning (ML) & Deep Learning (DL)
- ii. Brief historical overview
- iii. Applications for image analysis
- iv. When using DL and which tools are available?**

Why using Deep Learning?

When **classic image processing/analysis tools** are not efficient or do not exist for the task we want to perform (eg. PALM image prediction)

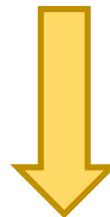


Need to have enough analyzed data to train the network



Need to label the data
in order to get
database large enough
for the training

Time consuming



Network are trained for
a specific set of data.
New type of data
means new training.

Not (always) flexible



Deep Learning needs
large computational
resources for image
analysis

Expensive

Which tools are available for Deep Learning?



Matlab 2018 version and later

Toolkits dedicated for ANN and deep learning.



Python 3 – open source

For DL, the open-source **TensorFlow** library from Google is used.



imJoy – open source

Plugin developed by the MIT and freely available on GitHub.



Deep Learning with Python



Open source distribution of **Python**, compatible with Windows, MacOS and Linux.

Spyder and **Jupyter Notebook** are two python environments available with Anaconda and compatible with **TensorFlow**

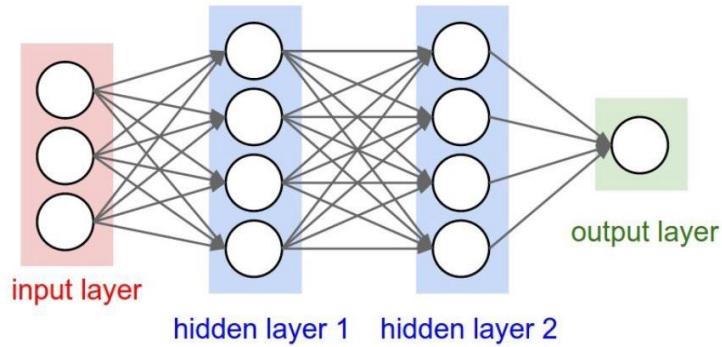


Outline

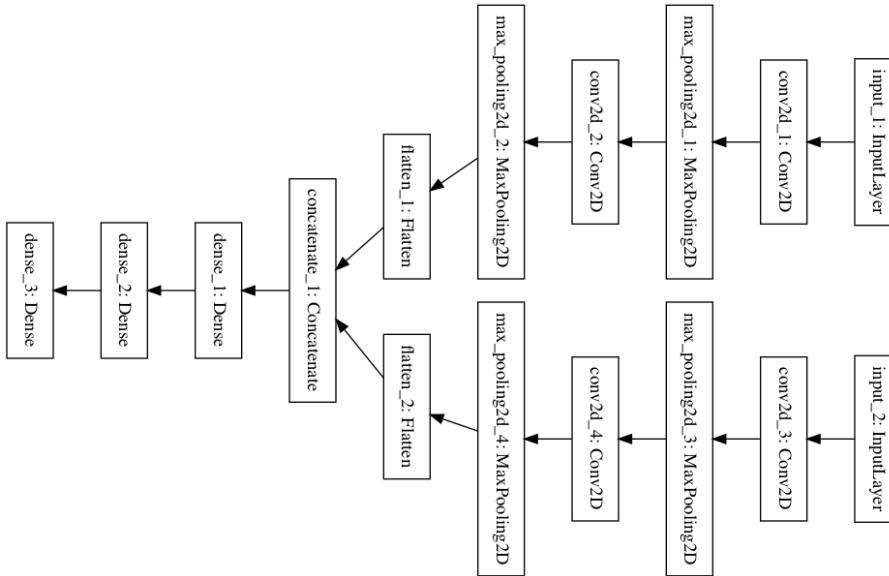
- I. Introduction**
- II. Starting with Deep Learning and Artificial Neuron Network**
 - i. A bit of vocabulary before starting**
 - ii. Case 1 : single neuron network
 - iii. Case 2 : how to construct a densely connected network
 - iv. Case 3 : multi-class clusterization
 - v. Summary

Sequential vs. non-sequential models

Sequential model

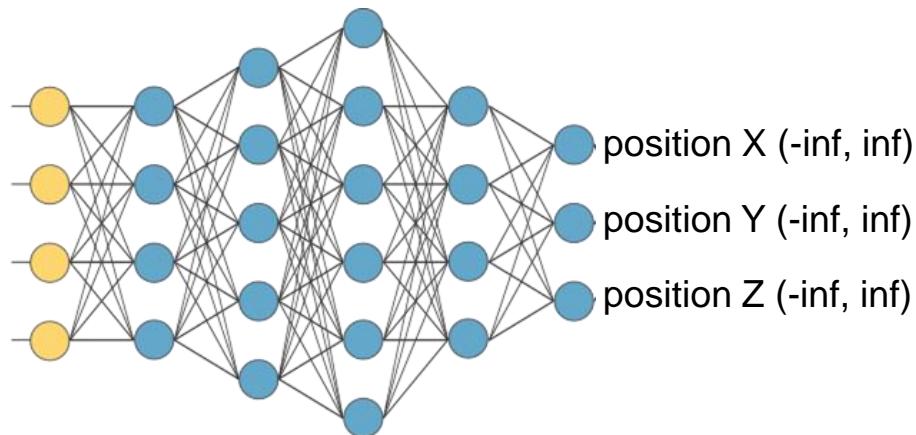
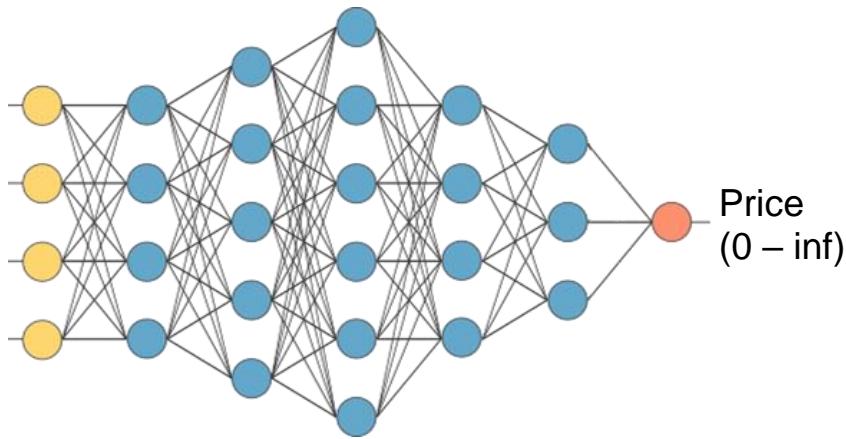


Non-sequential model

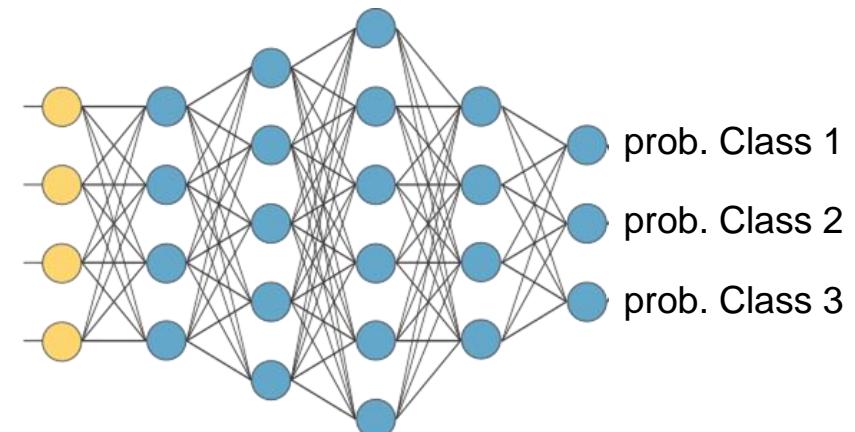
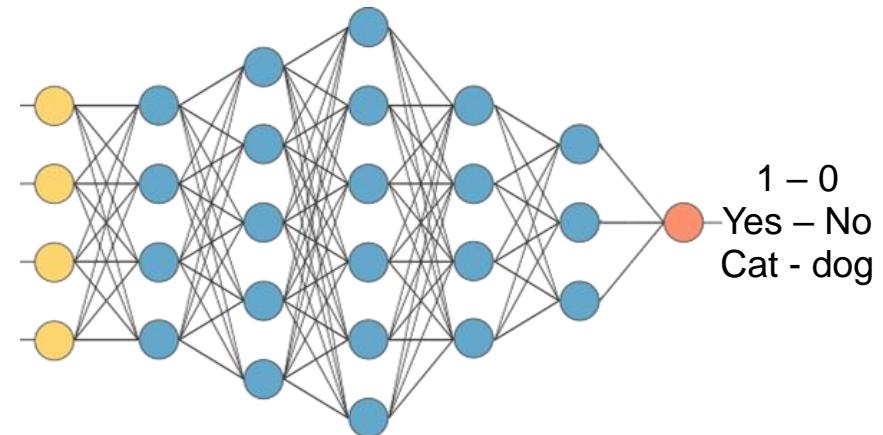


Regression vs. Classification

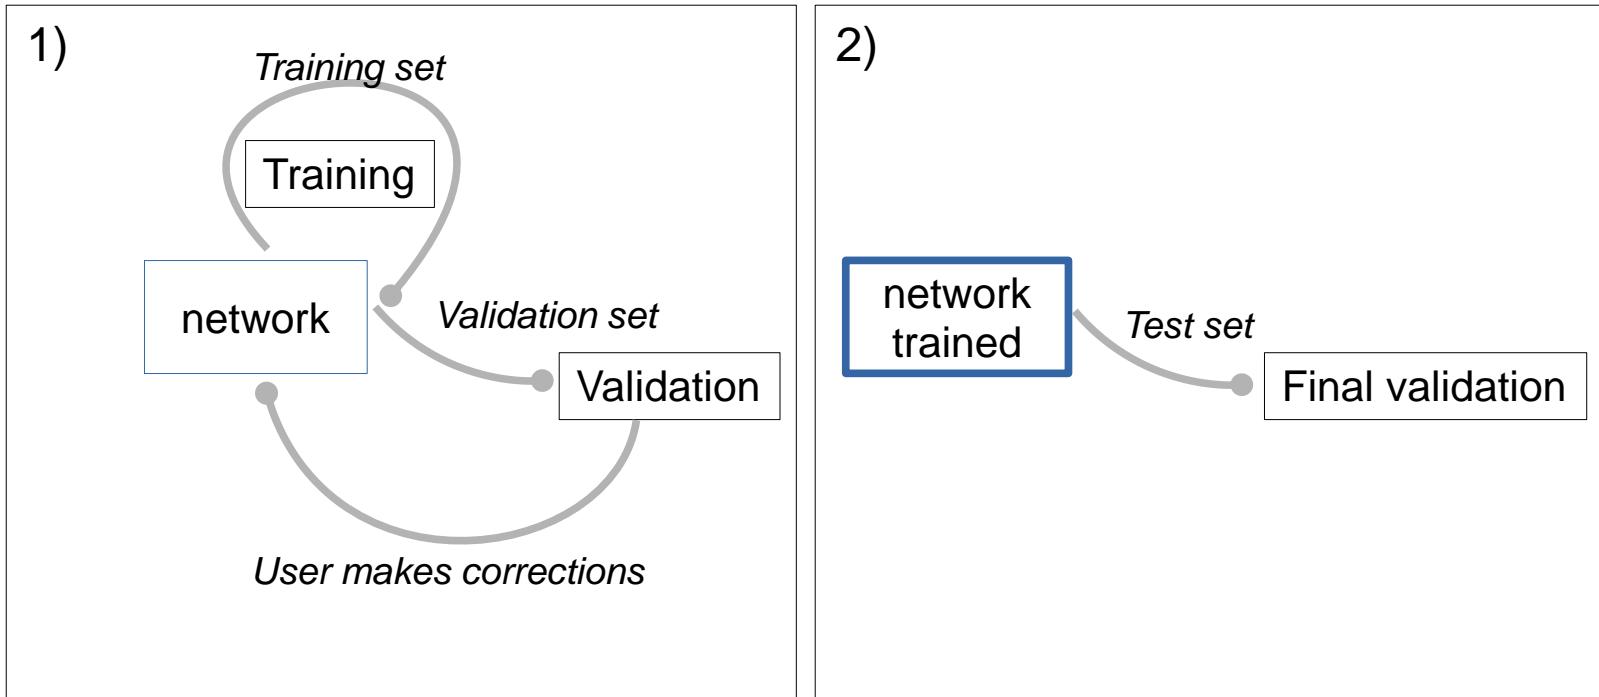
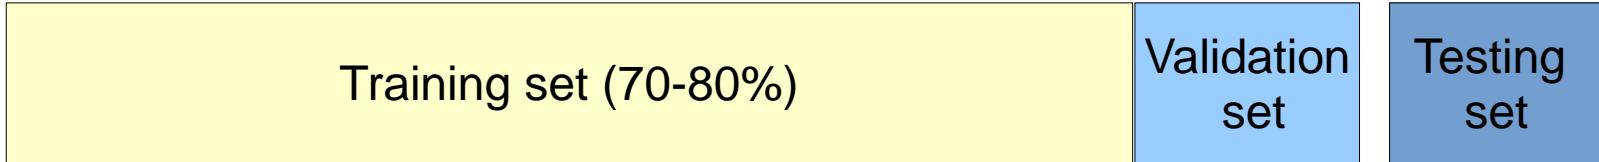
Regression : output is one or more real numbers



Classification : output is the probability that input belong to one or more classes



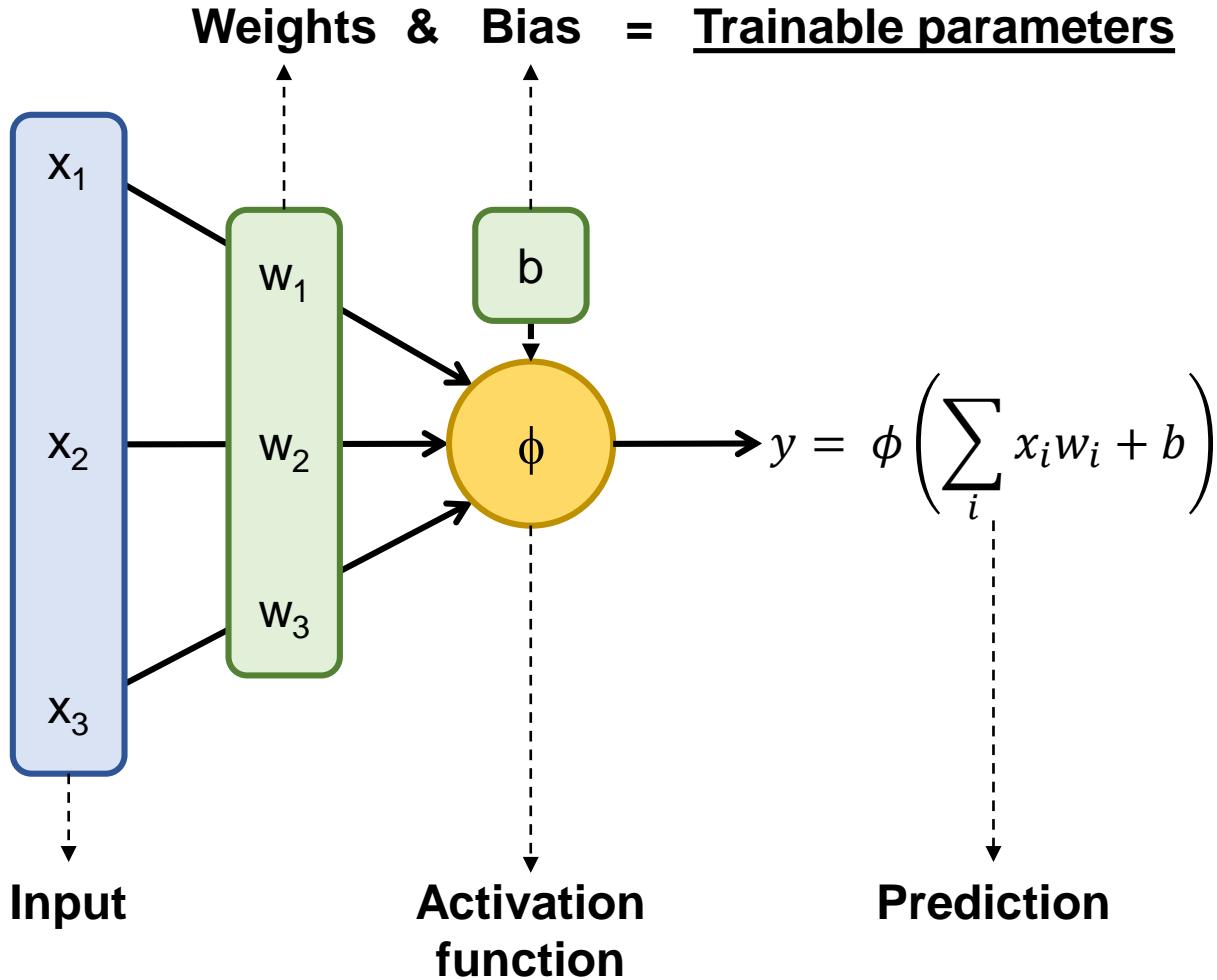
Training, Testing and validation sets



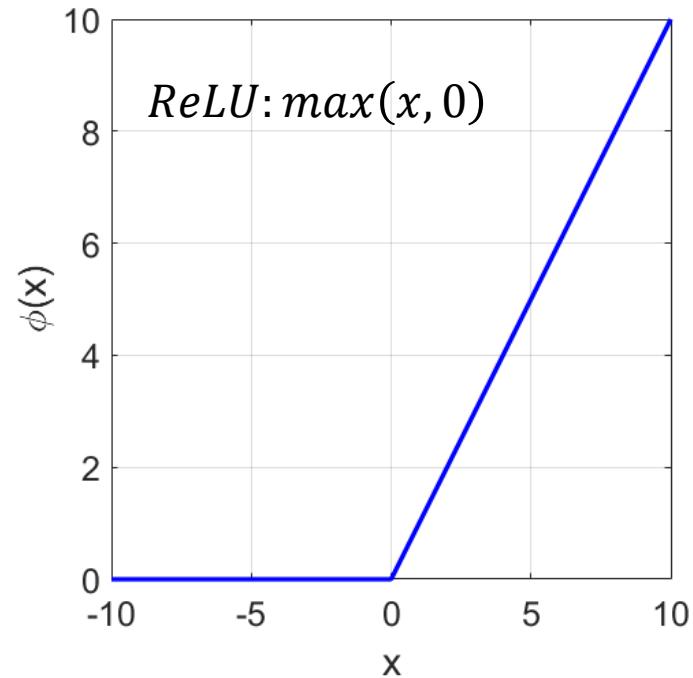
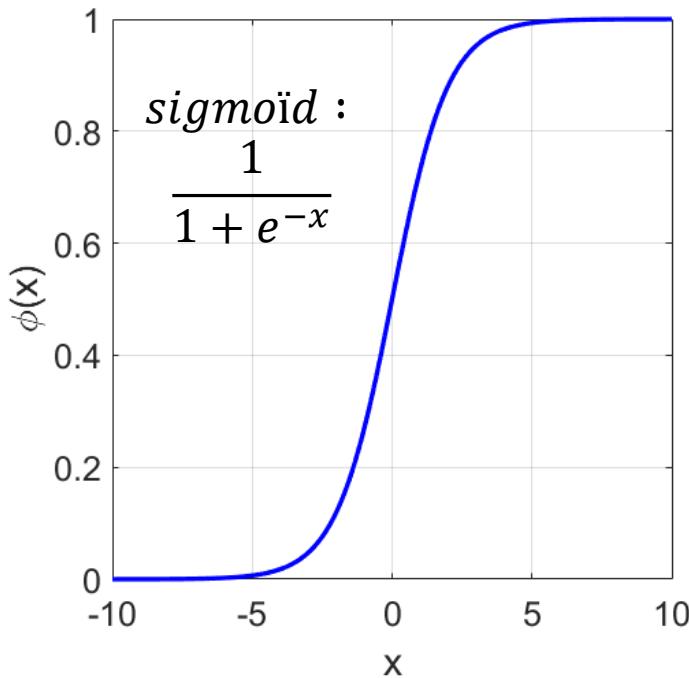
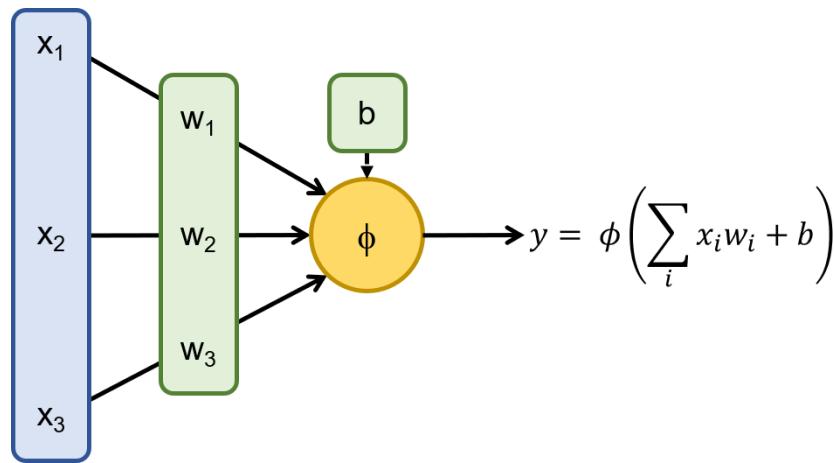
Outline

- I. Introduction**
- II. Starting with Deep Learning and Artificial Neuron Network**
 - i. A bit of vocabulary before starting
 - ii. Case 1 : single neuron network**
 - iii. Case 2 : how to construct a densely connected network
 - iv. Case 3 : multi-class clusterization
 - v. Summary

Definition of an artificial neuron



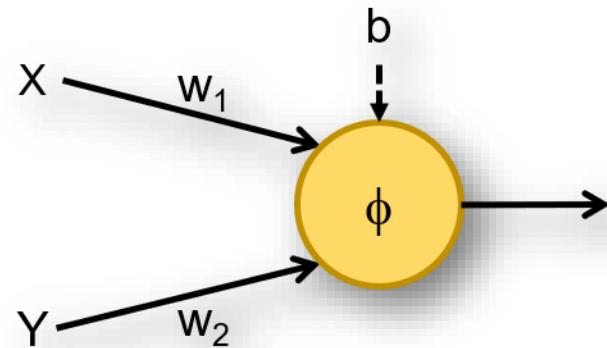
Most common activation functions



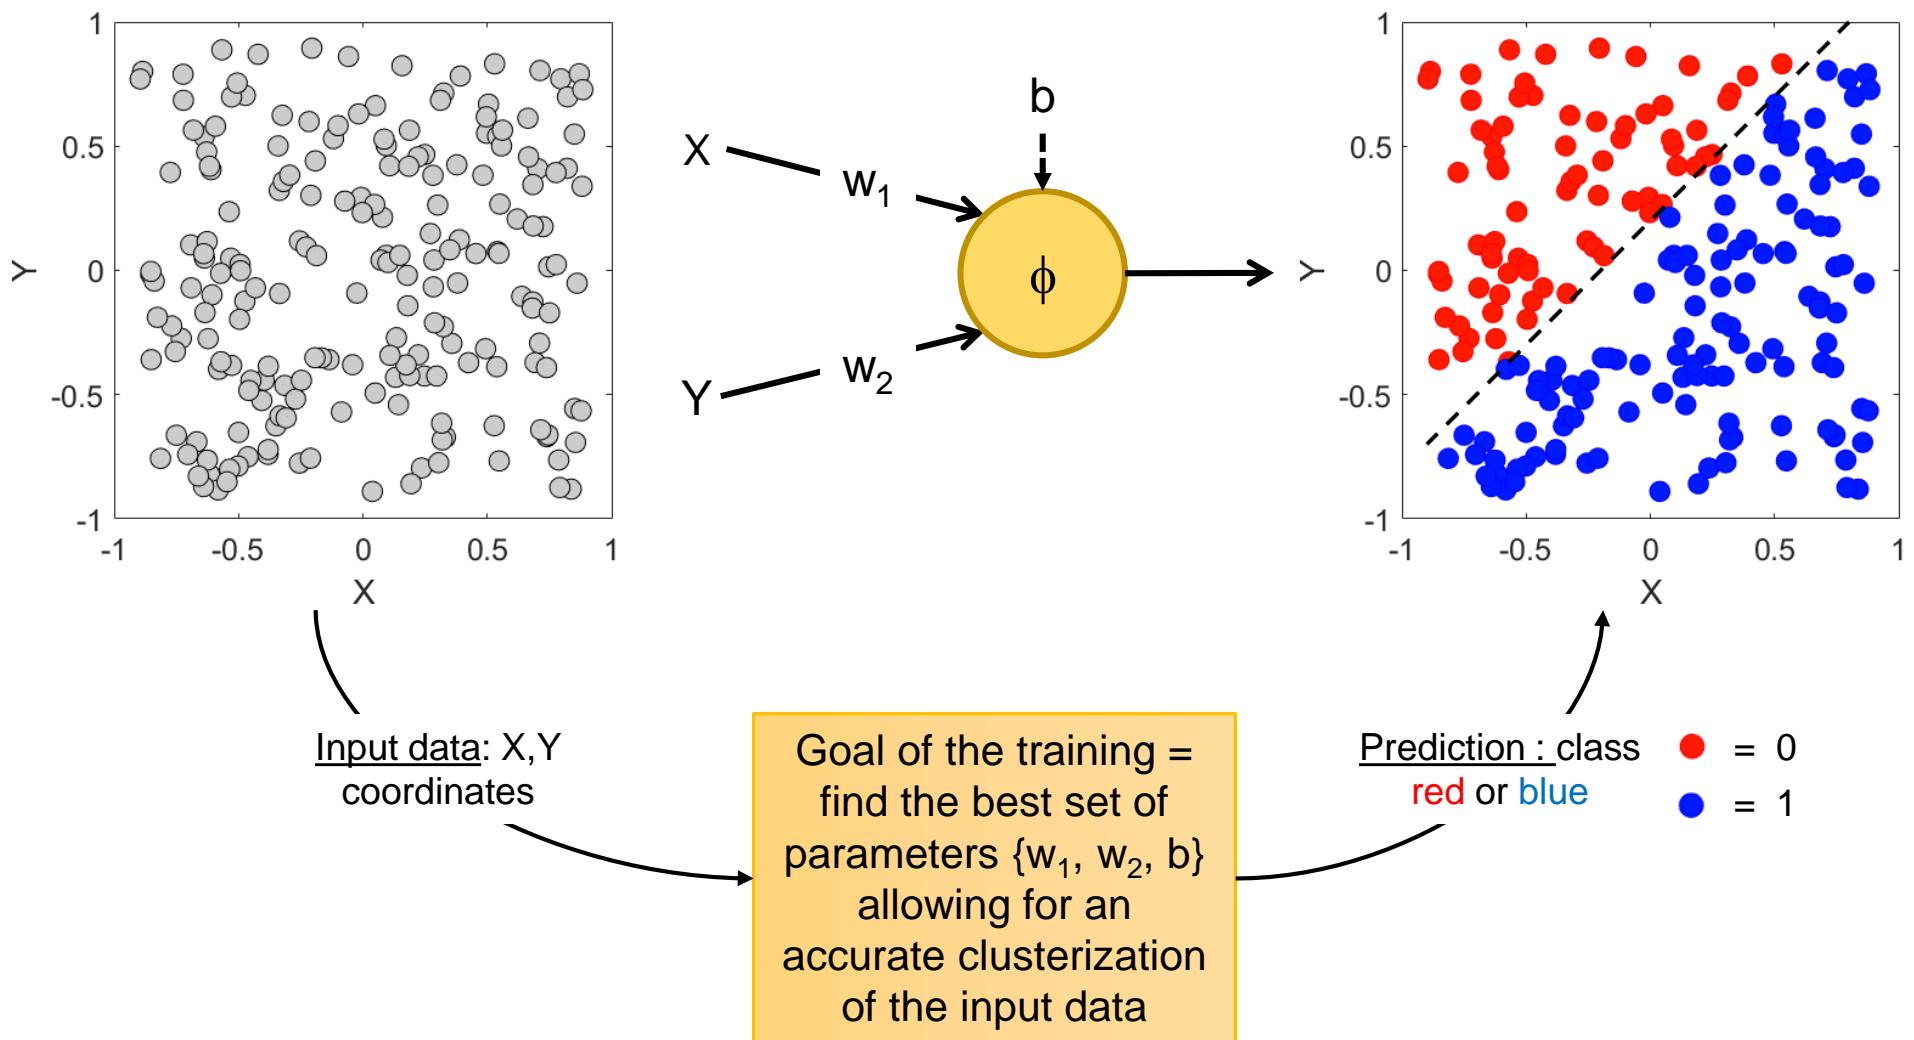
Example 1 : train a simple classifier with one neuron

Example n°1 : Clusterization_linearly_separated.ipynb

1. Understand the principle of the training
2. Train the classifier and test its accuracy
3. First step with Keras/TensorFlow



Example 1 : A one neuron classifier



Definition of the classifier with Keras

1- Definition of the network architecture

```
from keras import models
from keras import layers

model = models.Sequential()
model.add(layers.Dense(1, activation='sigmoid', input_shape=(2,)))
```

2- Definition of the training options

```
model.compile(optimizer='sgd',
              loss='binary_crossentropy',
              metrics=['accuracy'])
```

3- Training

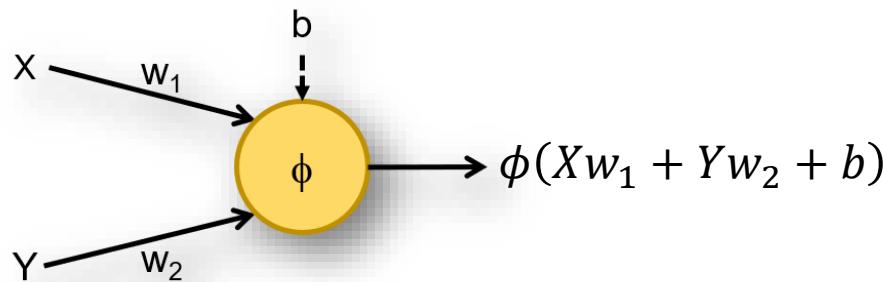
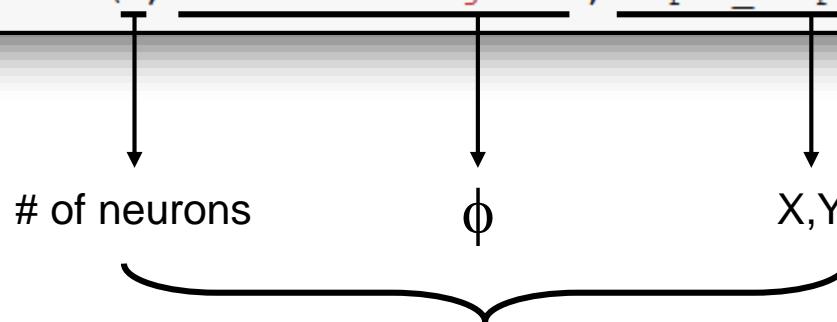
```
history = model.fit(Training_data,
                     Training_label,
                     epochs = 50,
                     validation_data = (Validation_data, Validation_label))
```

Definition of the classifier with Keras

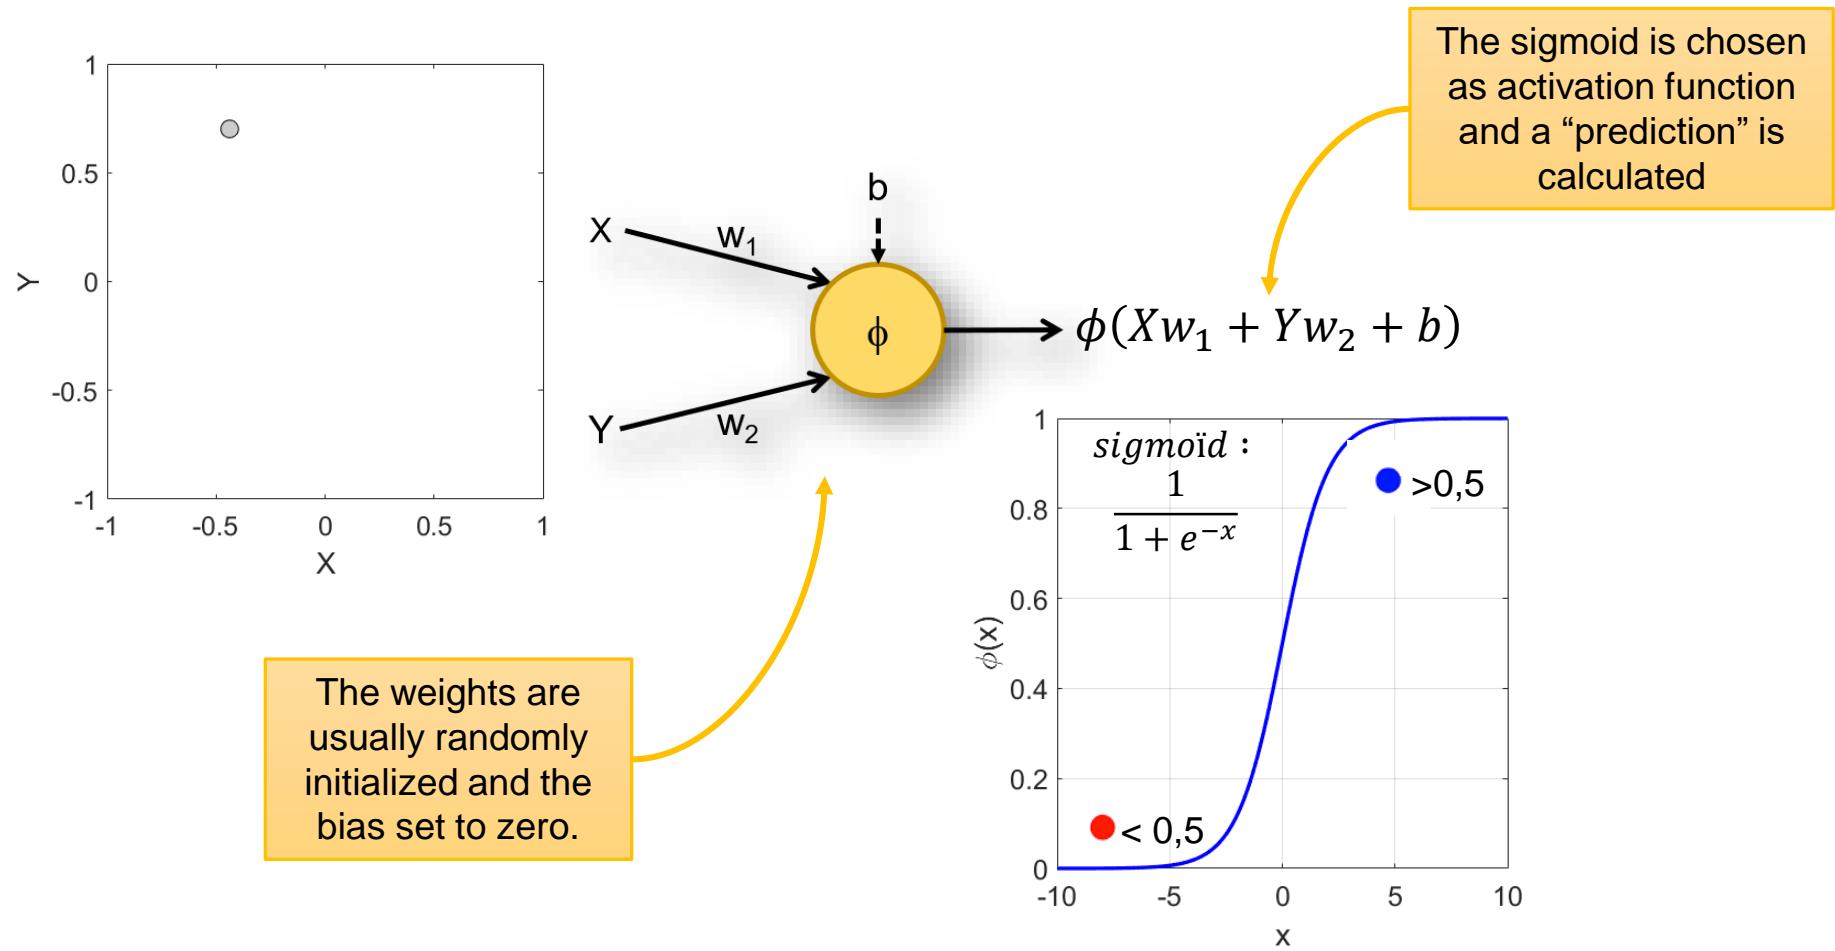
1- Definition of the network architecture

```
from keras import models
from keras import layers

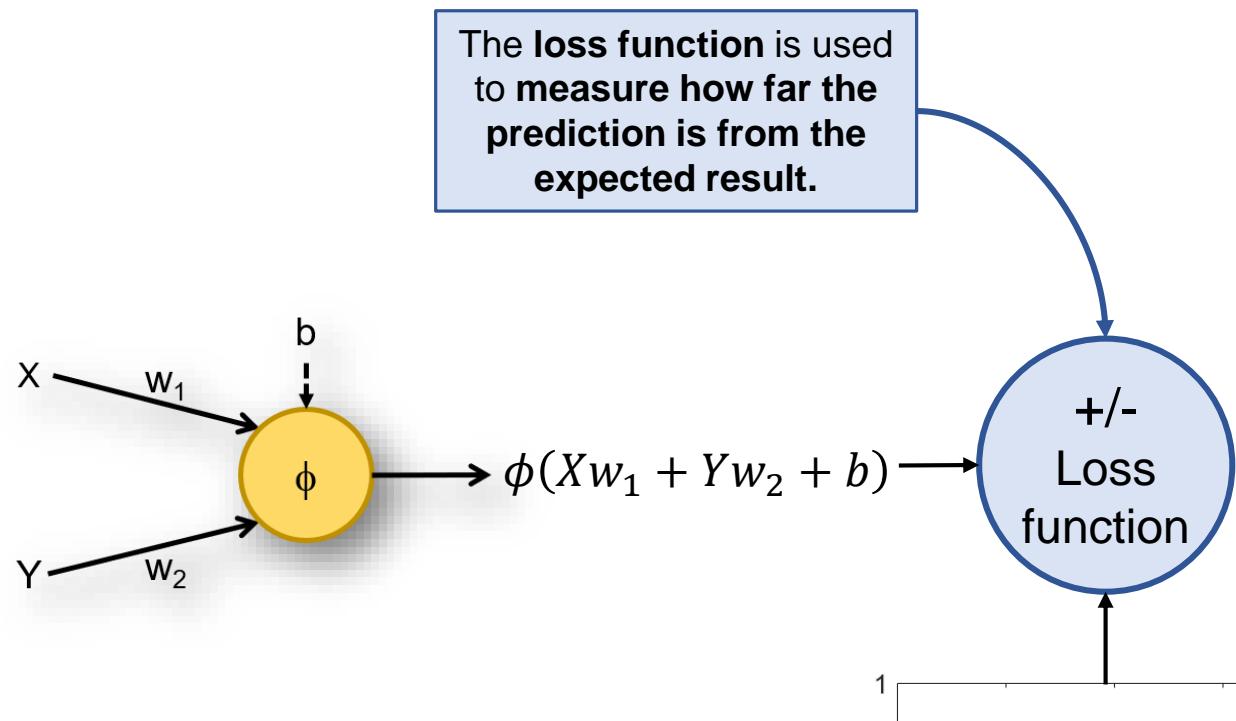
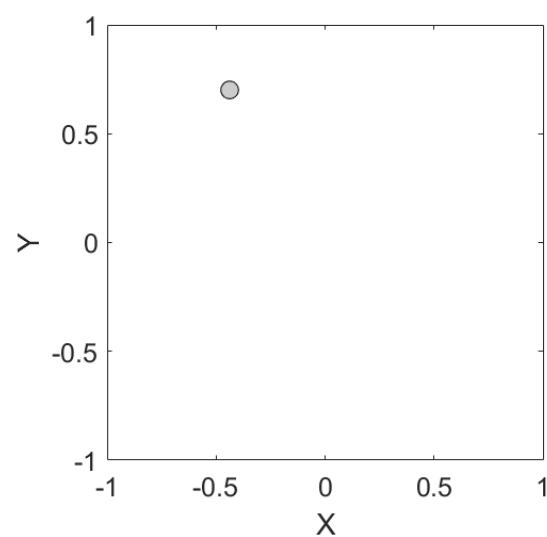
model = models.Sequential()
model.add(layers.Dense(1, activation='sigmoid', input_shape=(2,)))
```



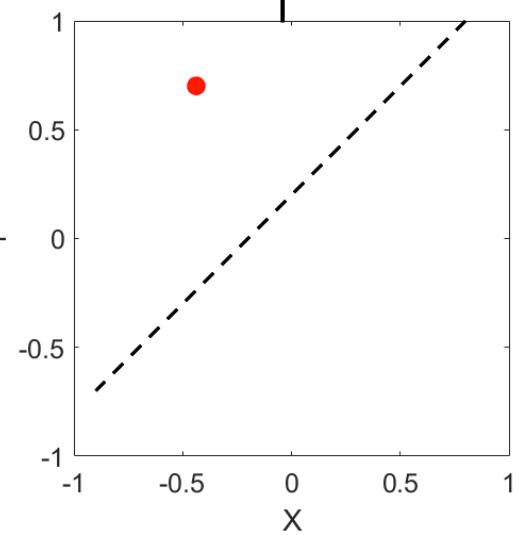
Training : step by step explanation



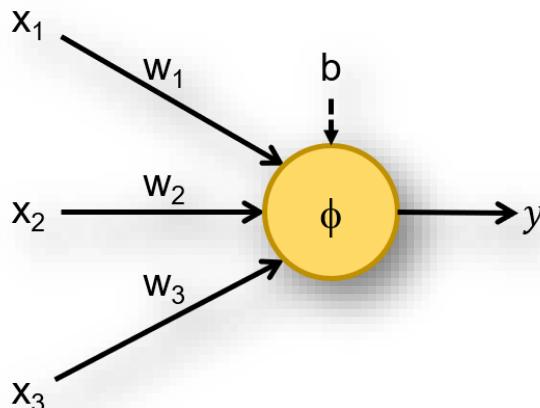
Training : step by step explanation



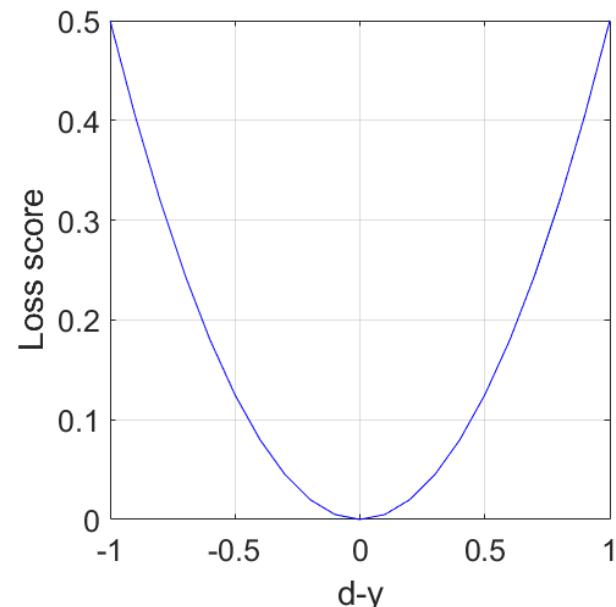
The calculated output from the neuron is compared to the expected result.



What is the loss function?

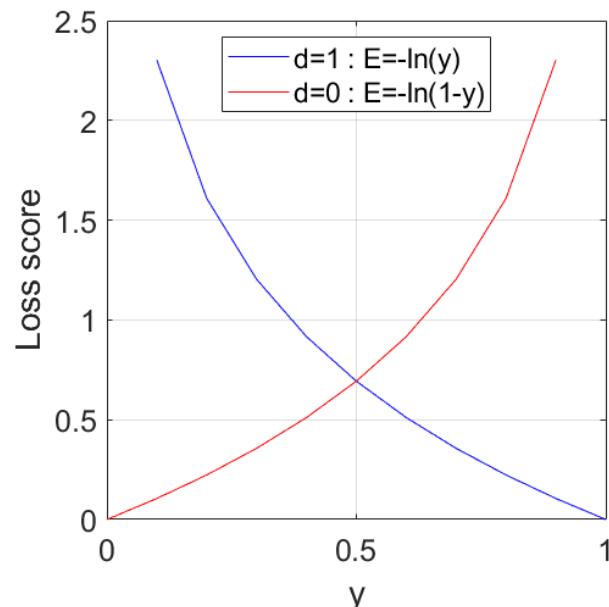


$$J = \frac{1}{2} (d_i - y_i)^2$$



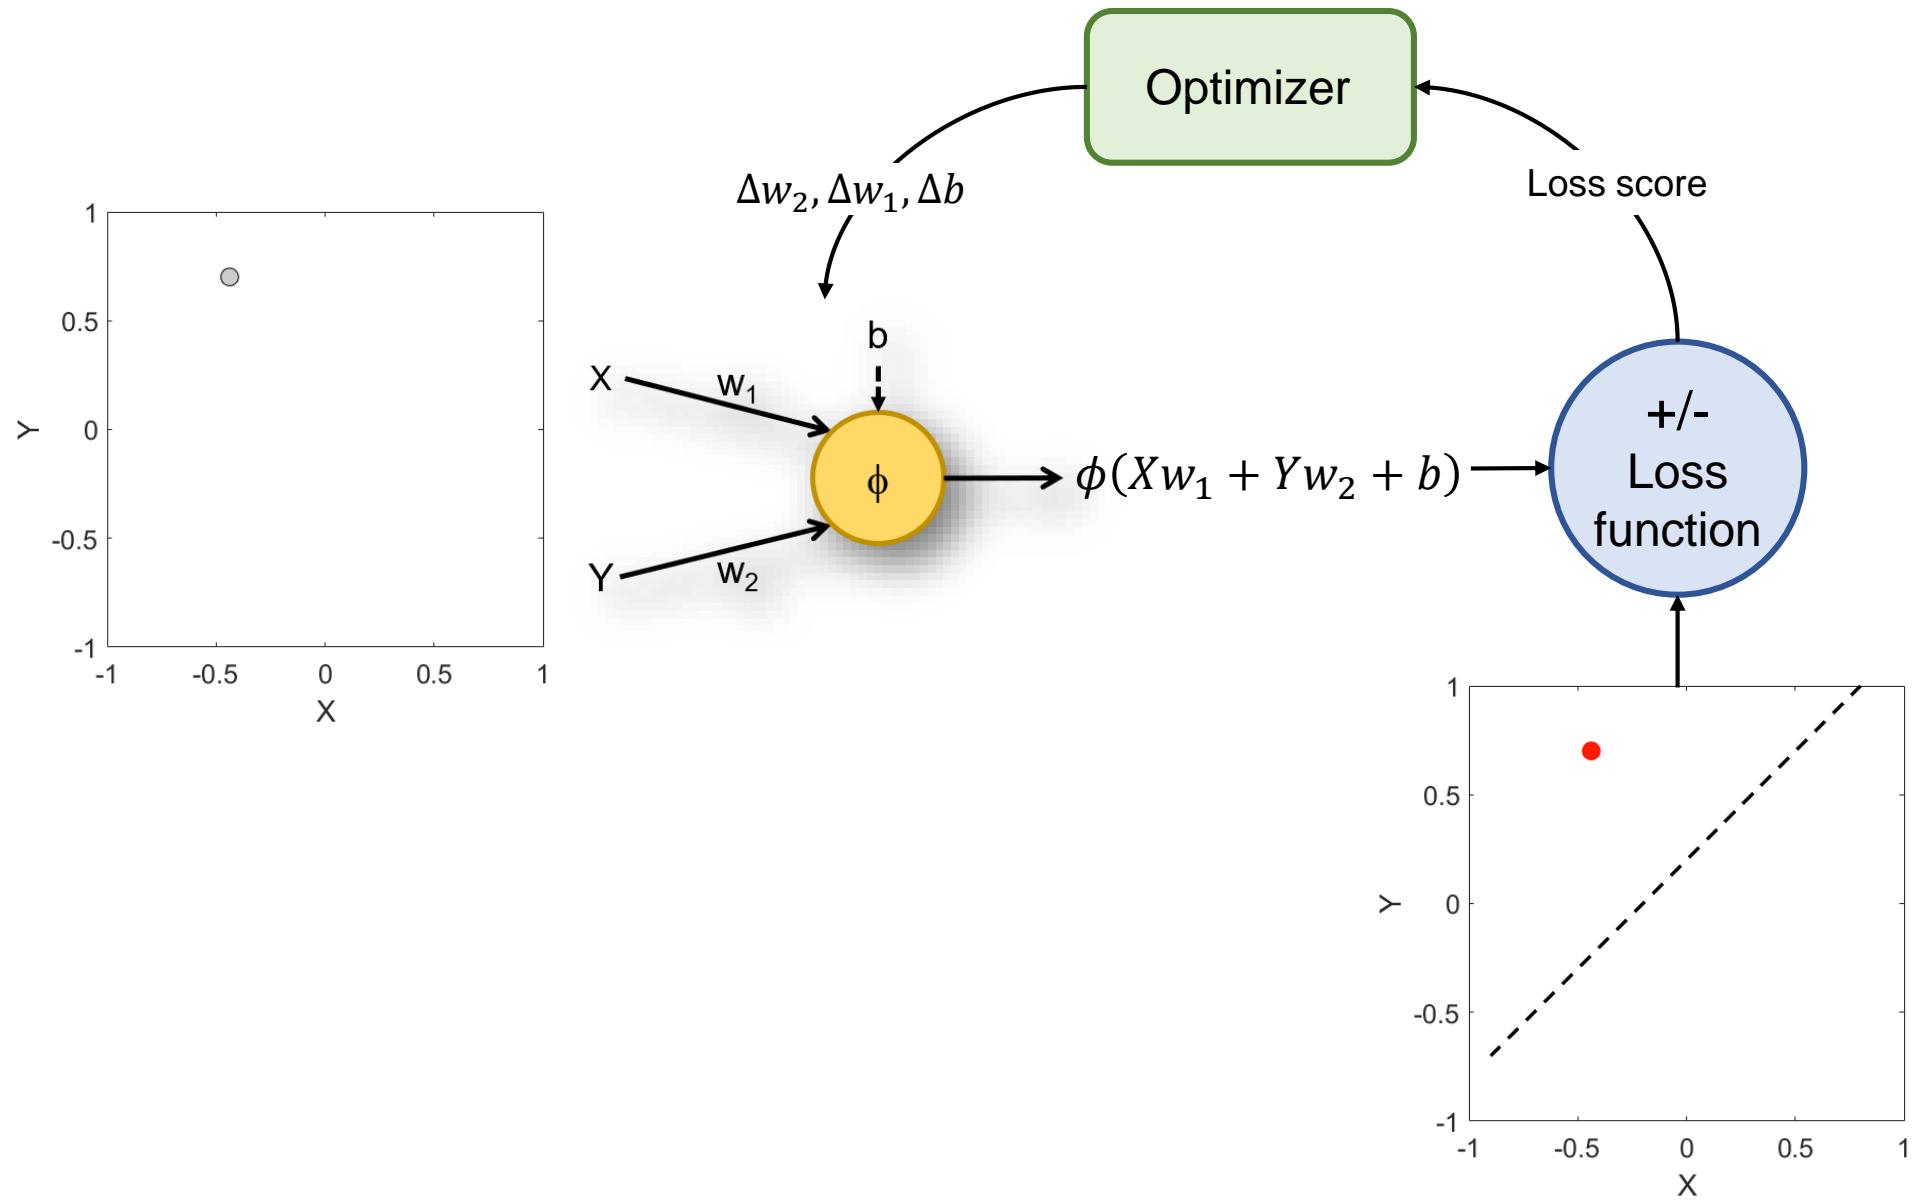
Squared error function, mainly used for regression problems.

$$J = \{-d_i \ln(y_i) - (1 - d_i) \ln(1 - y_i)\}$$

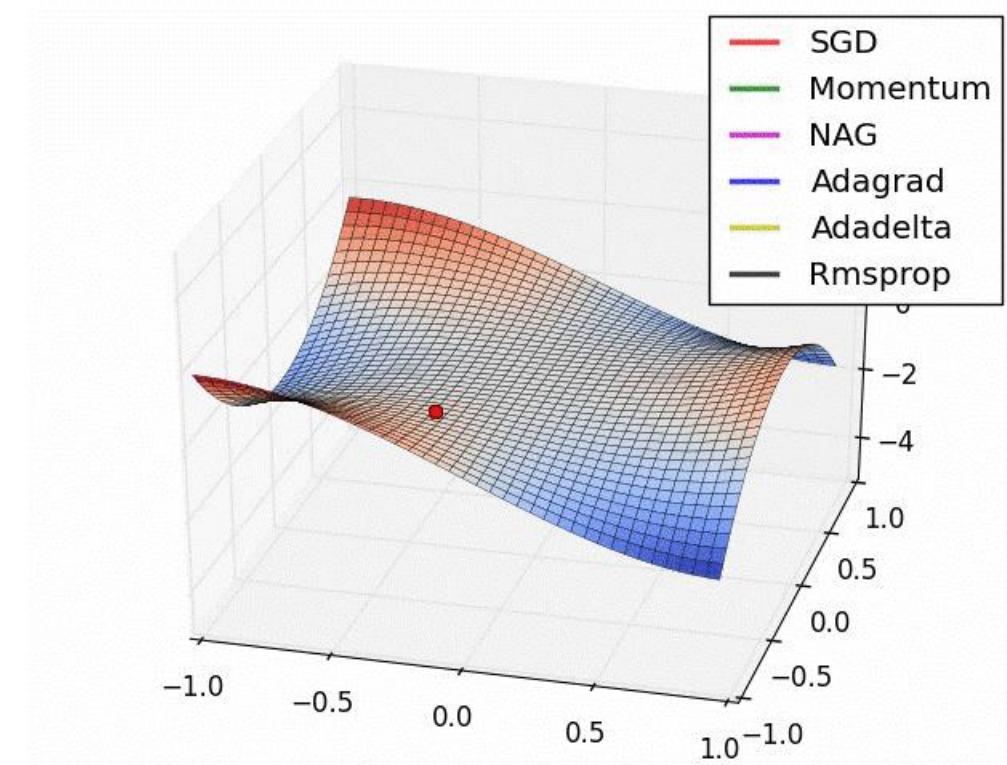
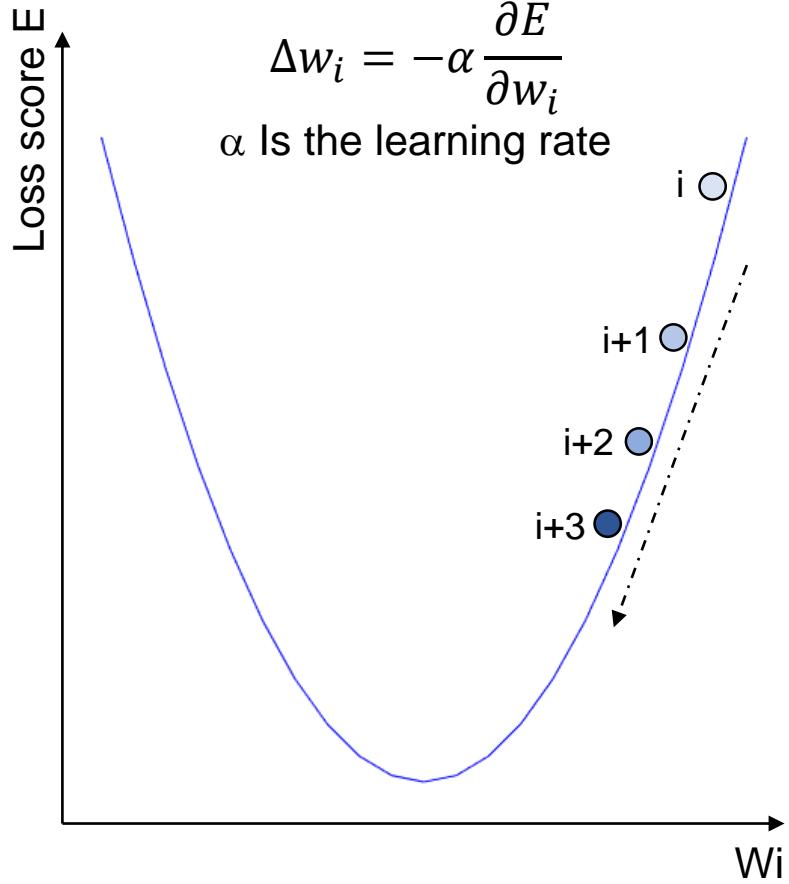


Binary crossed entropy, used in classification problems.

Training : step by step explanation



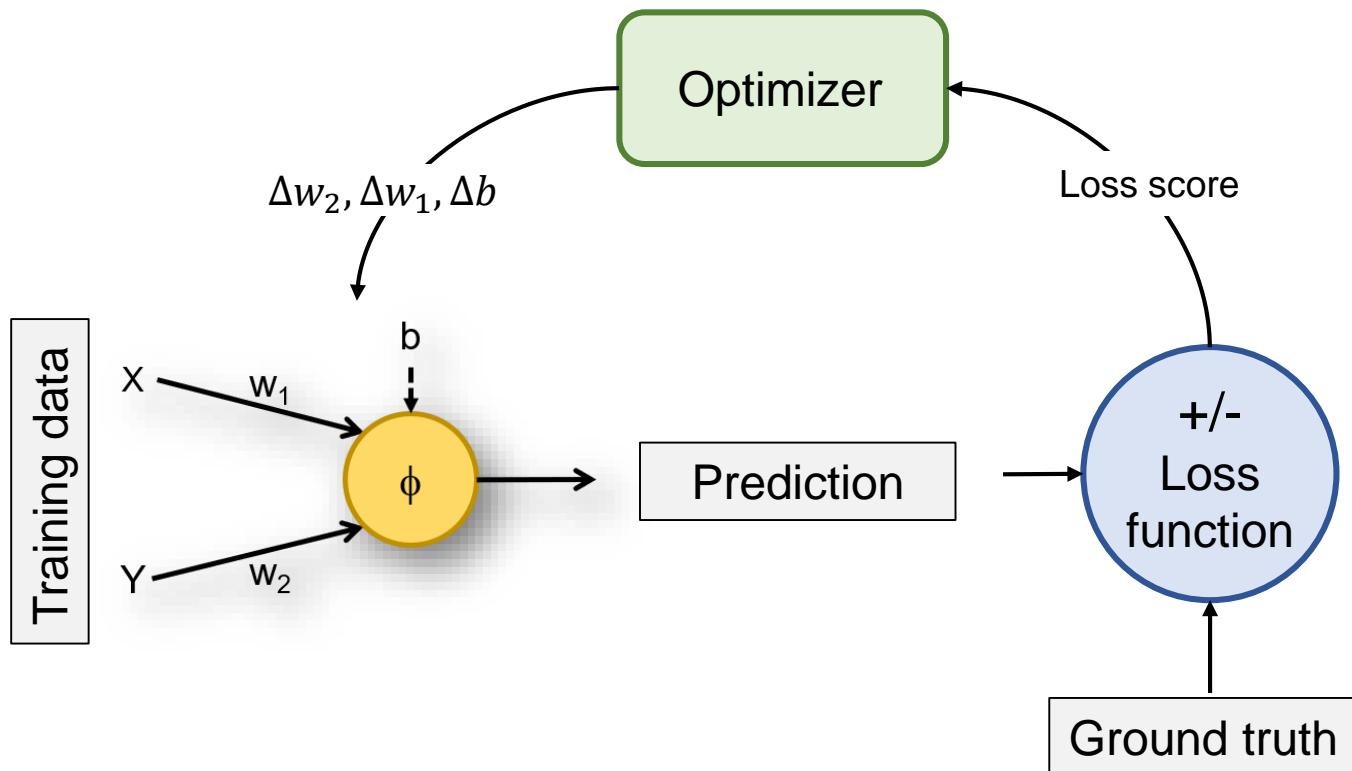
Optimizer : single gradient descent SGD



Training : step by step explanation

2- Definition of the training options

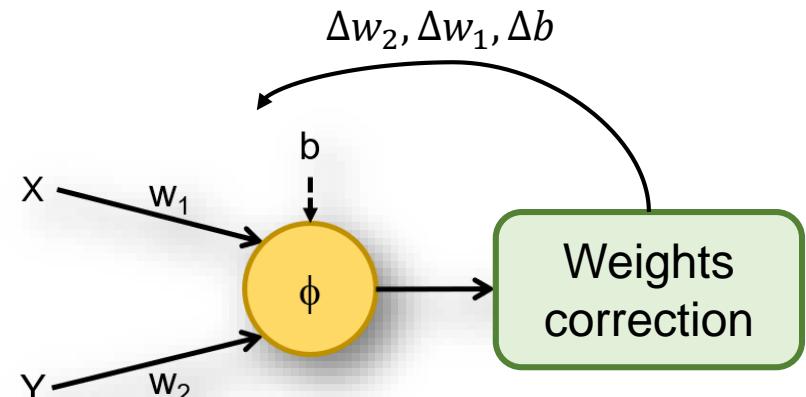
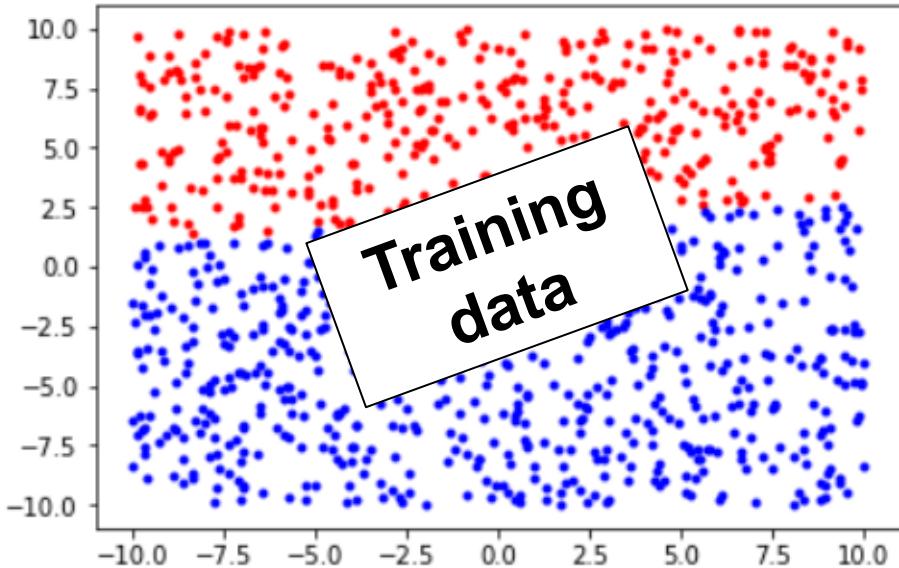
```
model.compile(optimizer='sgd',
              loss='binary_crossentropy',
              metrics=['accuracy'])
```



Training results

3- Training

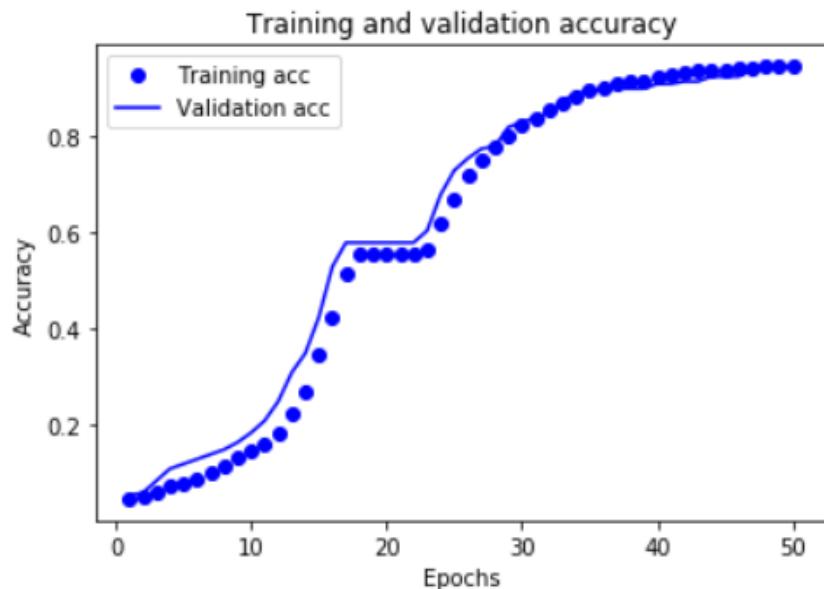
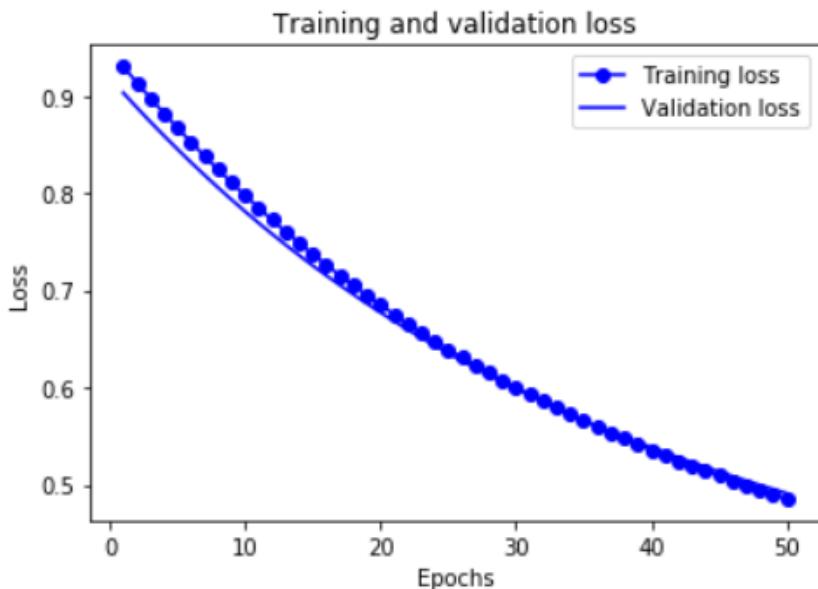
```
history = model.fit(Training_data,  
                    Training_label,  
                    epochs = 20,  
                    validation_data = (Validation_data, Validation_label))
```



Training results

3- Training

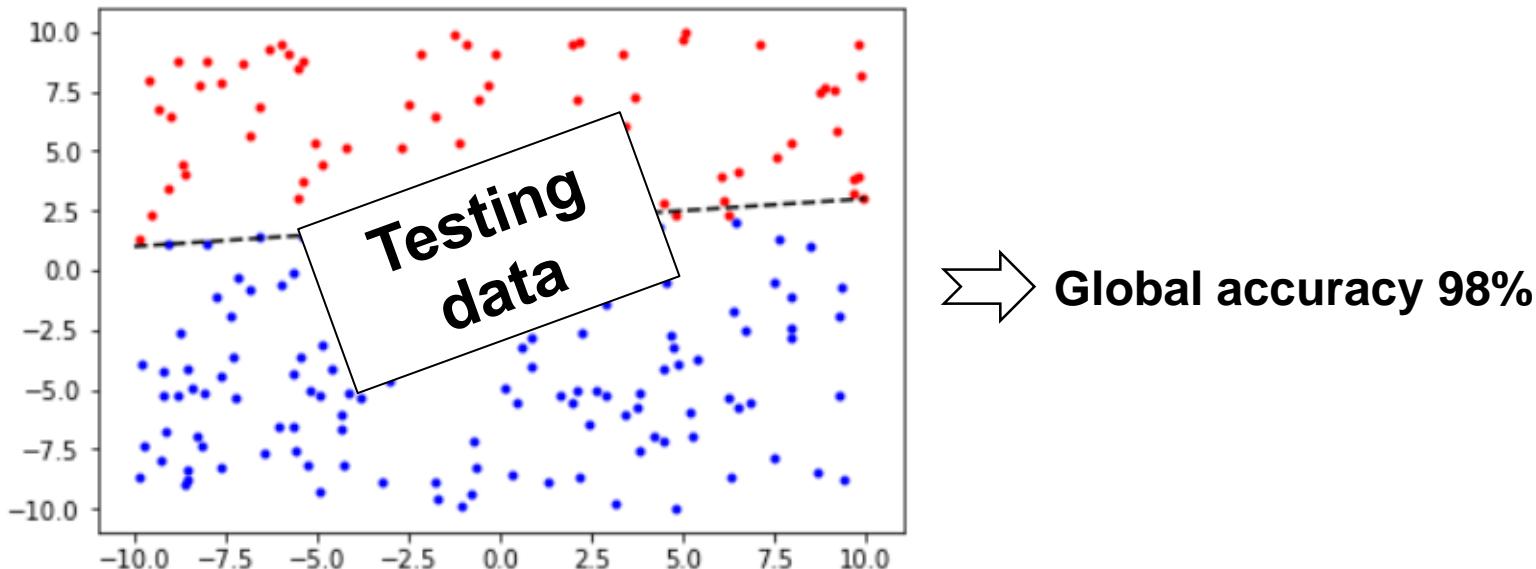
```
history = model.fit(Training_data,  
                    Training_label,  
                    epochs = 50,  
                    validation_data = (Validation_data, Validation_label))
```



Validation of the training using the testing set

Estimation of the model accuracy

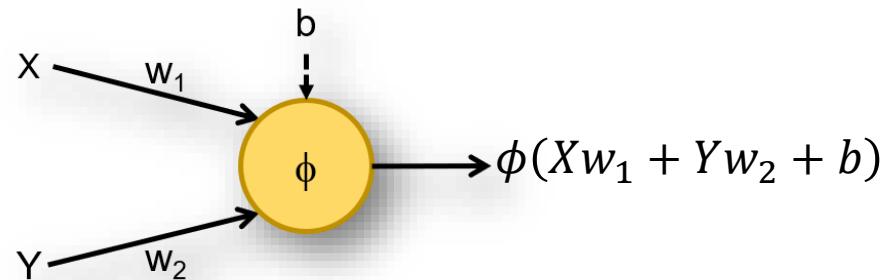
```
Predication_accuracy = model.evaluate(Testing_data, Testing_label)  
print(Predication_accuracy)
```



What did we learn so far?

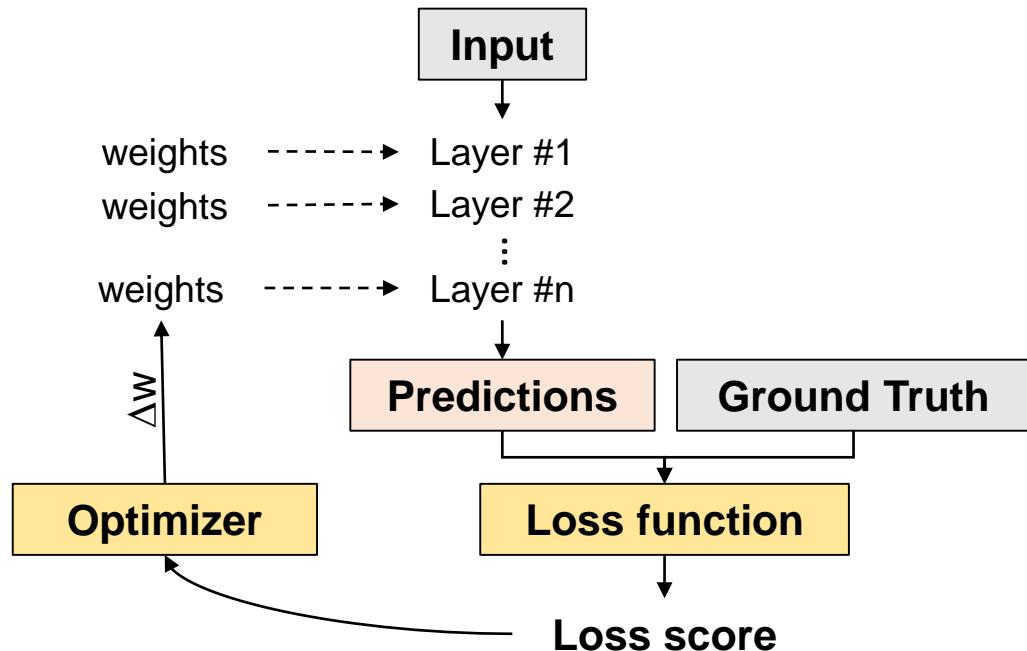
- How a single neuron works :

- Activation function
- Input / output
- Weights and bias



- How the training works :

- Loss function
- Optimizer
- Learning rate



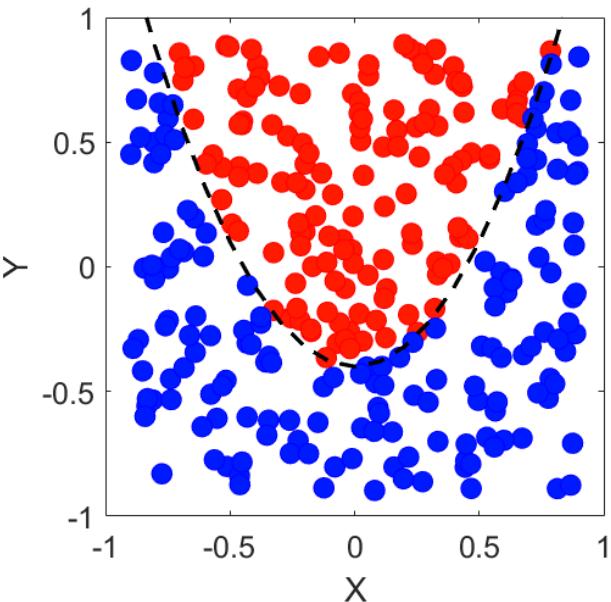
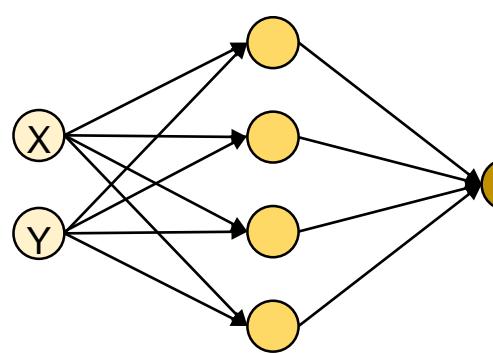
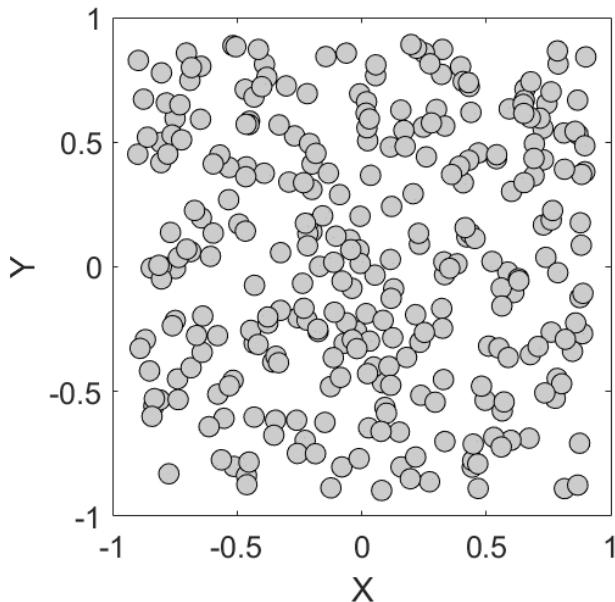
Outline

- I. Introduction
- II. Starting with Deep Learning and Artificial Neuron Network
 - i. A bit of vocabulary before starting
 - ii. Case 1 : single neuron network
 - iii. Case 2 : how to construct a densely connected network**
 - iv. Case 3 : multi-class clusterization
 - v. Summary

Example 2 : classify non-linearly separable data

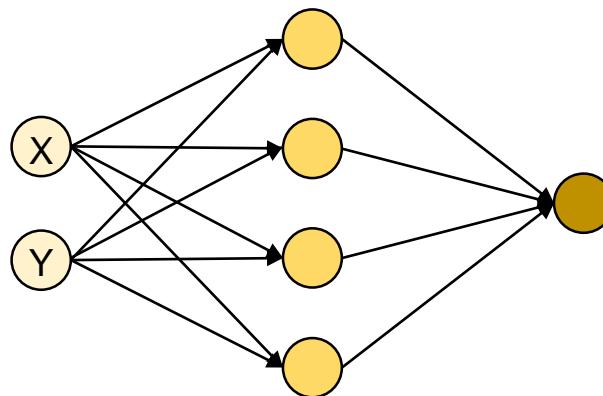
Example n°2 : Clusterization_not_linearly_separated_parbole

1. Observe the limitations of single-layer model
2. Find a simple architecture able to solve this classification problem



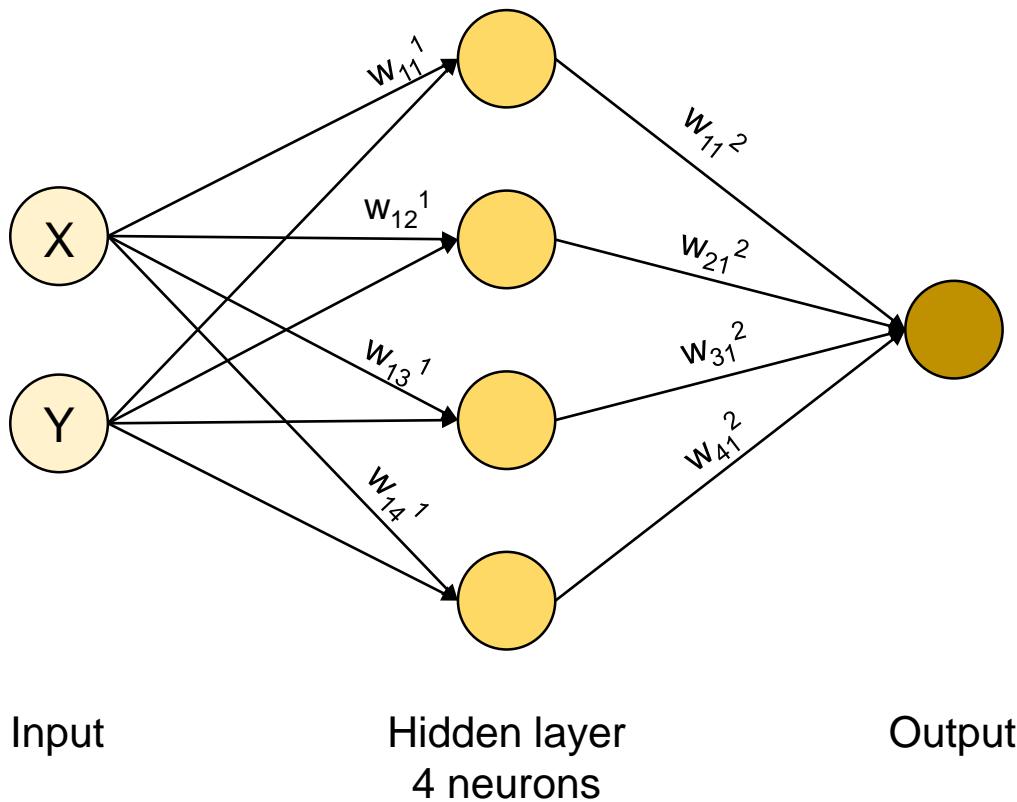
Example 2 : classify non-linearly separable data

- I. Try to train the classifier with one single neuron :
- II. Repeat the same procedure using more neurons to “fit” the model :



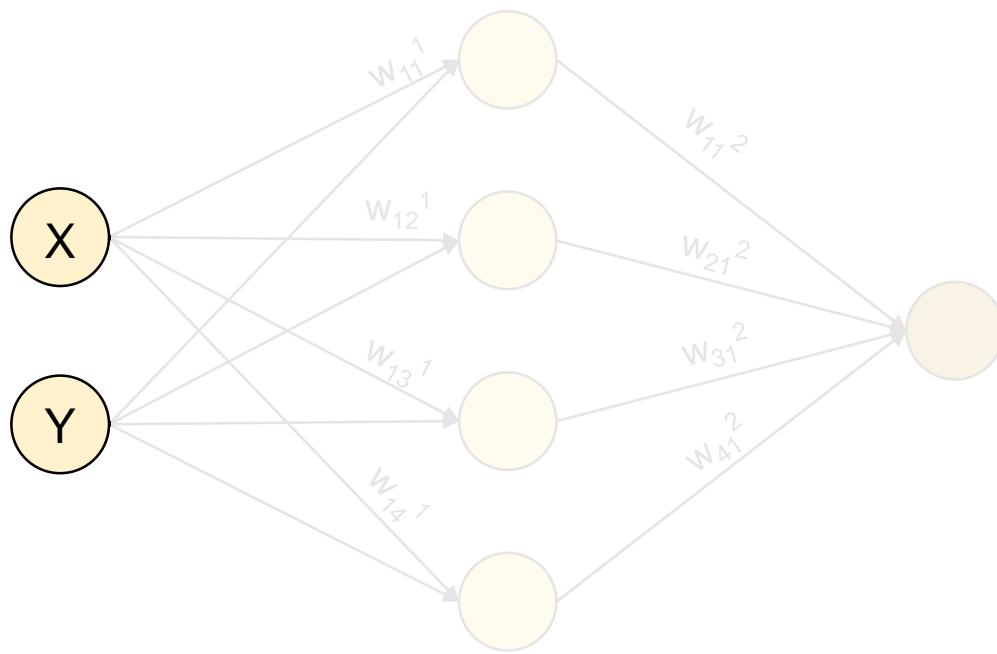
```
model = models.Sequential()  
model.add(layers.Dense(4, activation='sigmoid', input_shape=(2,)))  
model.add(layers.Dense(1, activation='sigmoid'))  
  
model.compile(optimizer = optimizers.sgd(lr=5e-2),  
              loss='binary_crossentropy',  
              metrics=['accuracy'])
```

Back-propagation algorithm : a quick glimpse



Total number of **trainable parameters** : $4*2 + 4 + 4 + 1 = 17$

Back-propagation algorithm : a quick glimpse



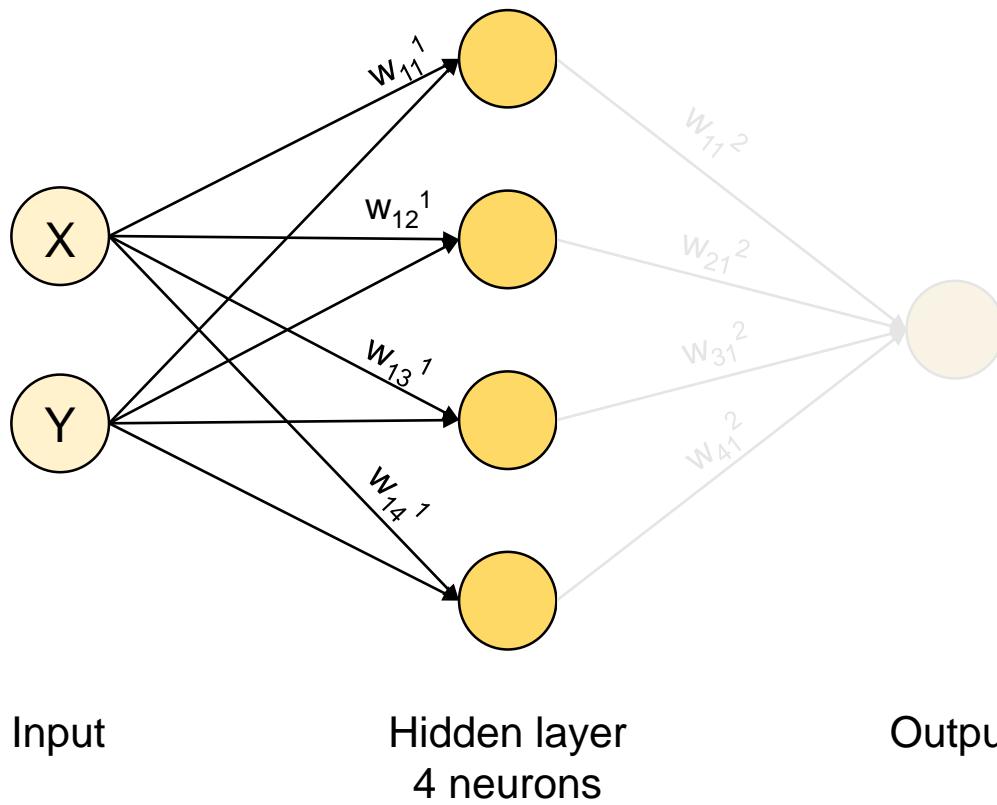
Input

Hidden layer
4 neurons

Output

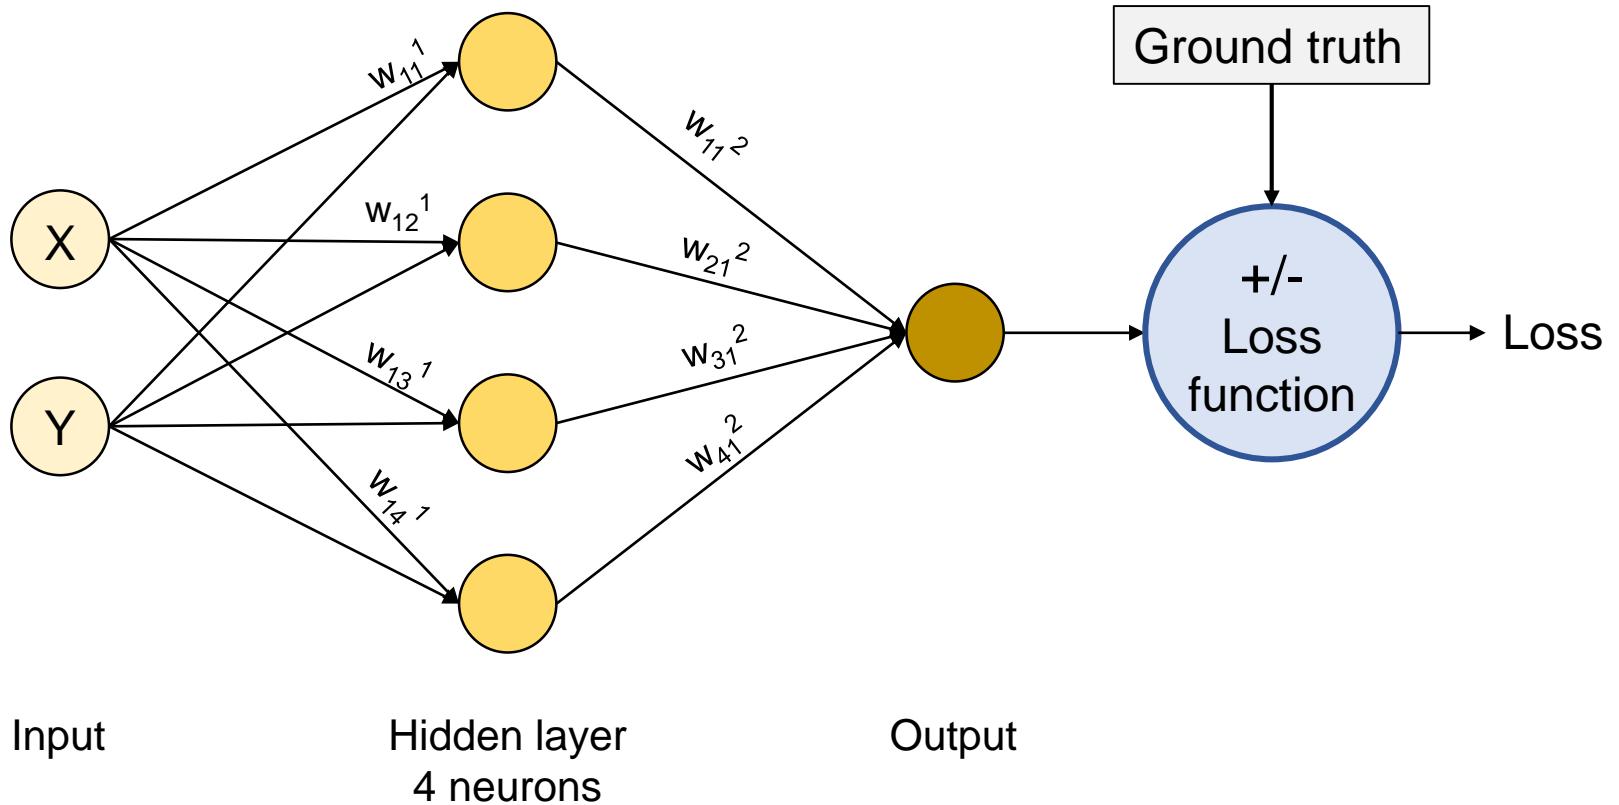
Forward pass

Back-propagation algorithm : a quick glimpse



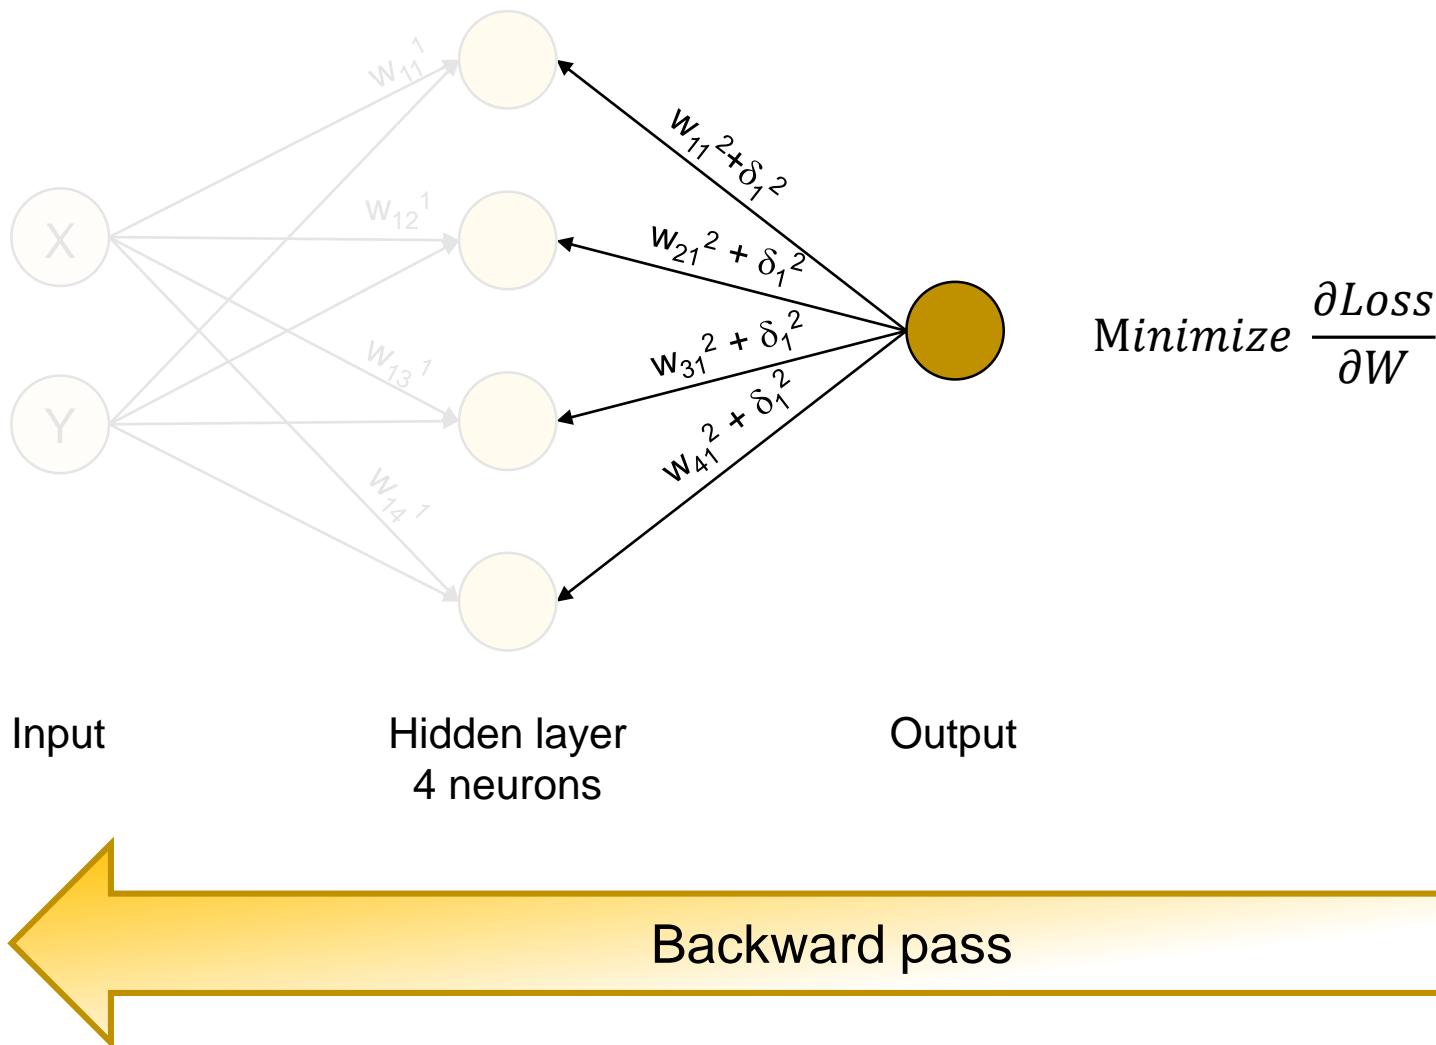
Forward pass

Back-propagation algorithm : a quick glimpse

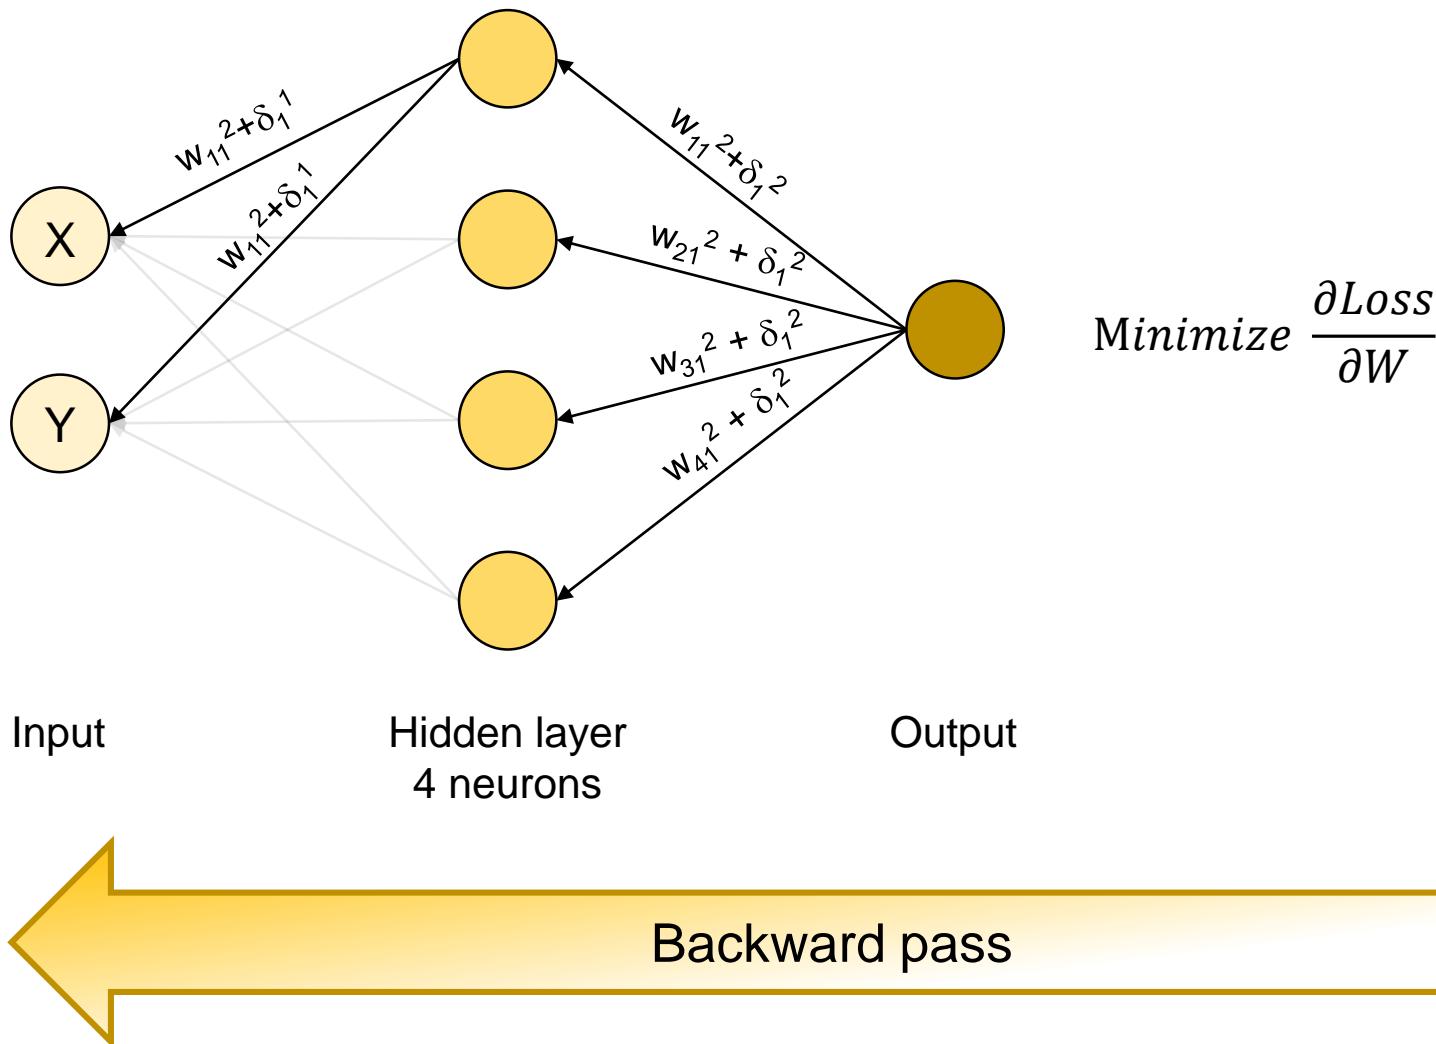


Forward pass

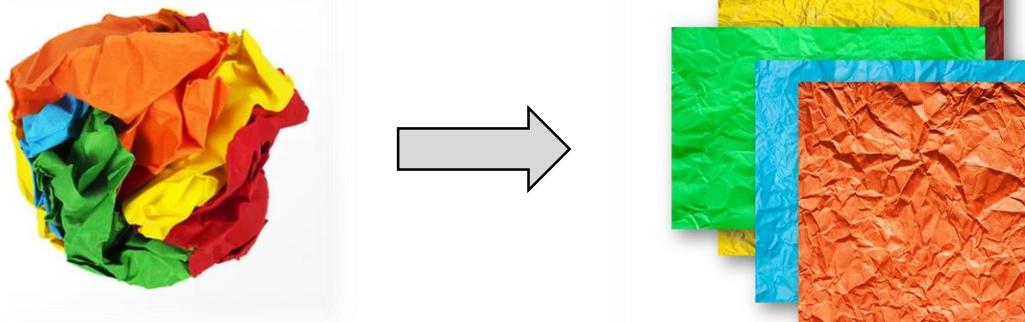
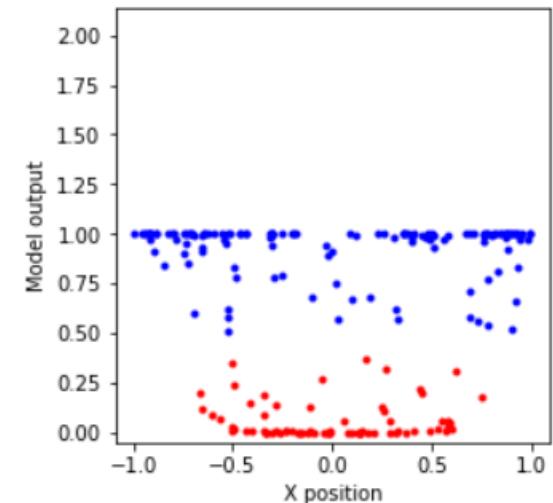
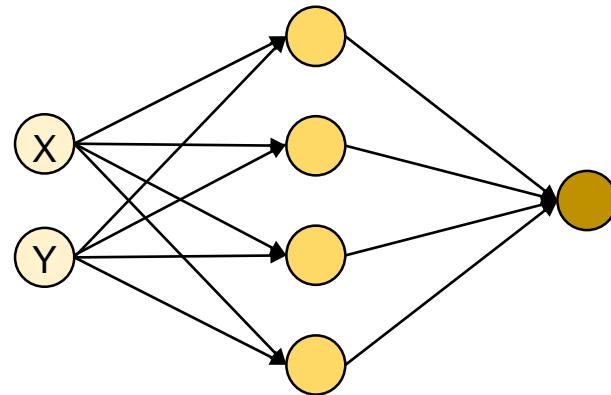
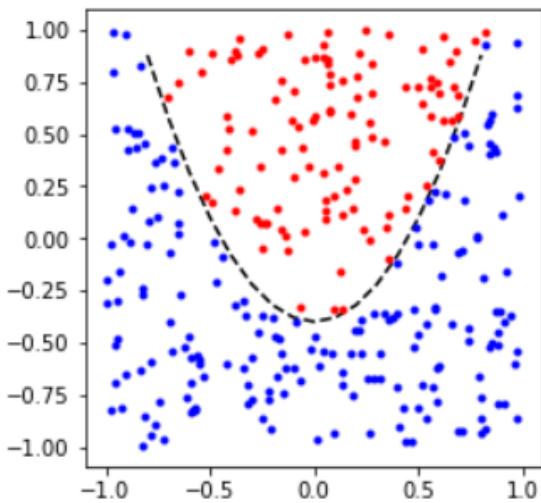
Back-propagation algorithm : a quick glimpse



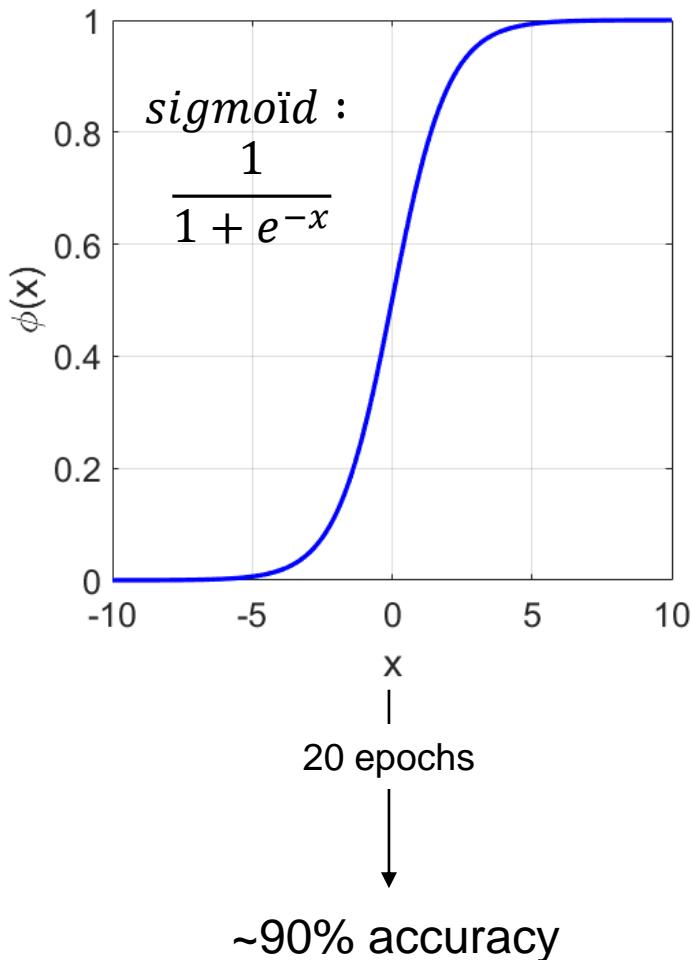
Back-propagation algorithm : a quick glimpse



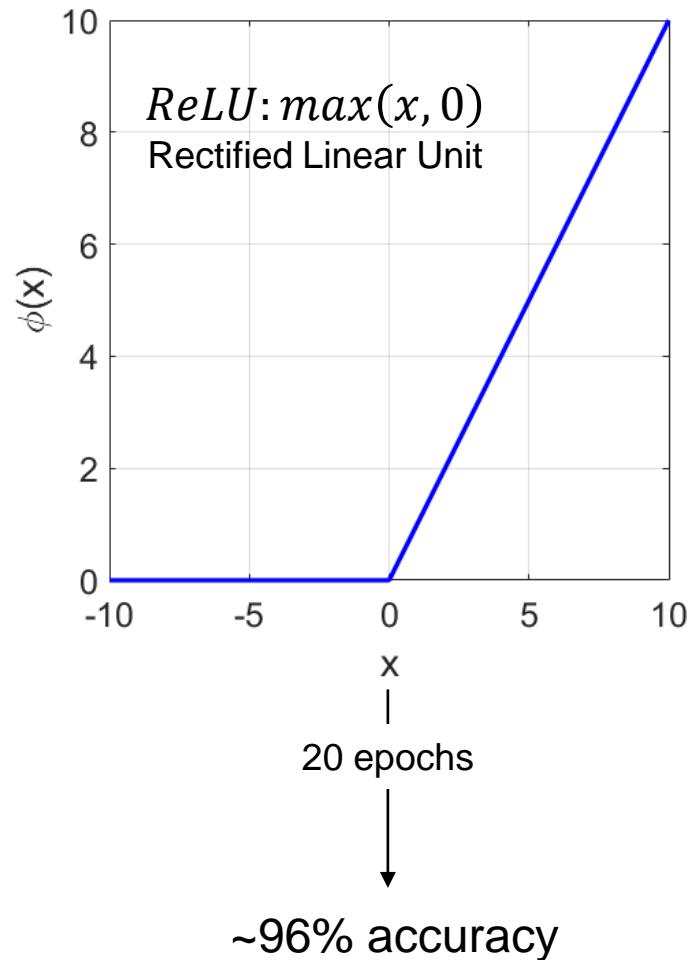
The network is learning a “better” representation



ReLU vs. Sigmoid



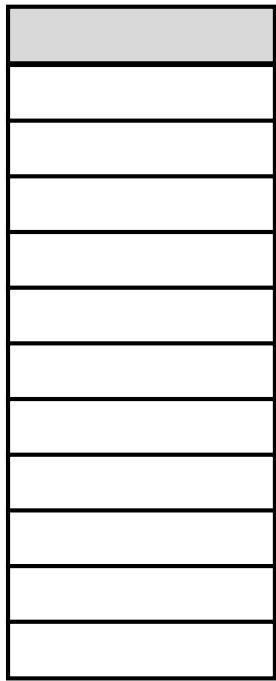
Gradient is saturating for large output values and $\exp()$ is a computer-expensive operation



Non-saturating gradient and very fast operation. However, negative values are discarded.

Batch processing

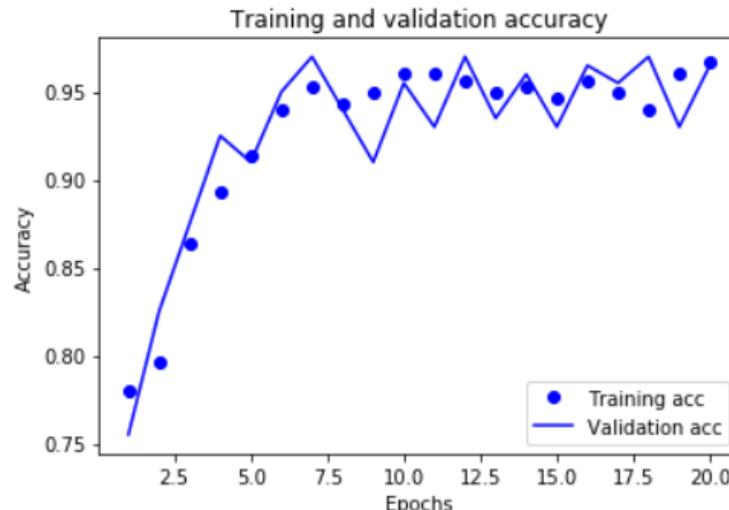
```
history = model.fit(Training_data,  
                    Training_label,  
                    epochs = 20,  
                    batch_size = 1,  
                    validation_data = (Validation_data, Validation_label))
```



For each training data
the weights are adjusted

$$\Delta w_i$$

The weights are
adjusted N times per
epoch

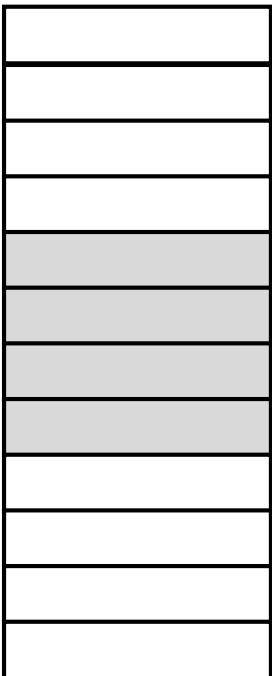


Training set = N

Batch processing

```
history = model.fit(Training_data,  
                    Training_label,  
                    epochs = 50,  
                    batch_size = 4  
                    validation_data = (Validation_data, Validation_label))
```

Better to use a power of 2



For each batch, the weights **are adjusted once** using the average

$$\frac{1}{batch_size} \sum \Delta w_i$$

The weights are adjusted **N/batch_size**

What did we learn so far?

- How to build a multi-layer network :

```
model = models.Sequential()
model.add(layers.Dense(5, activation='relu', input_shape=(2,)))
model.add(layers.Dense(5, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
```

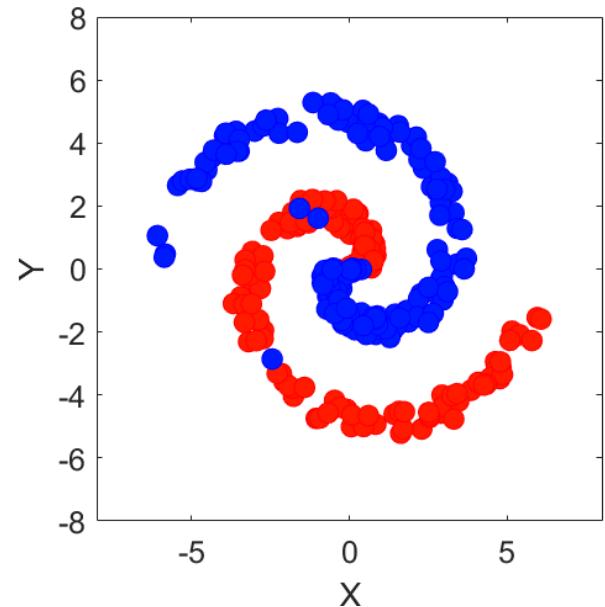
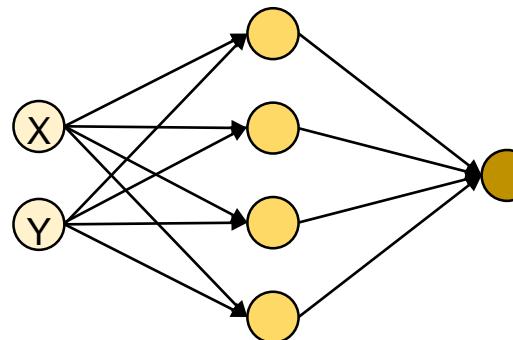
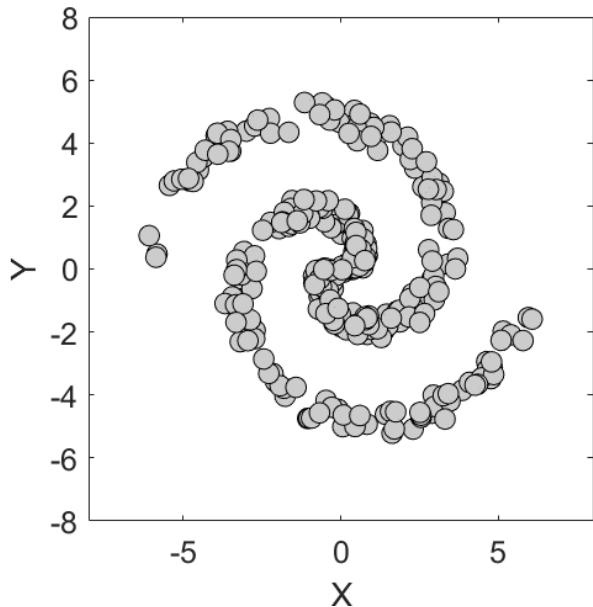
- The network learning capability increases with the **number of layers & neurons** → **complex data will require deeper network**
- But be careful to avoid **overfitting**
- For the training :
 - The **relu activation function** is more effective and allow the network to learn faster
 - Using **mini-batch** allows for a faster and more stable training

<https://playground.tensorflow.org/>

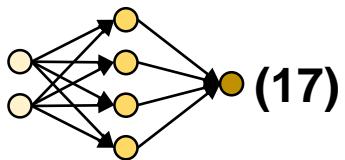
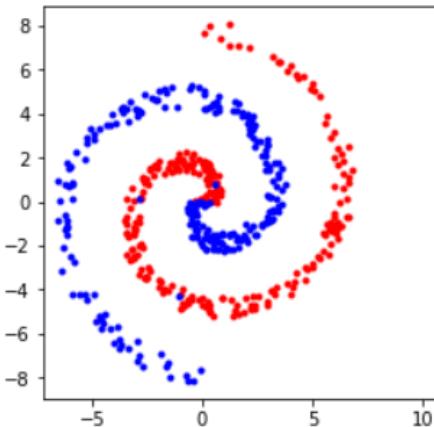
<https://cs230.stanford.edu/lecture/>

Challenge : build your own ANN to solve a problem

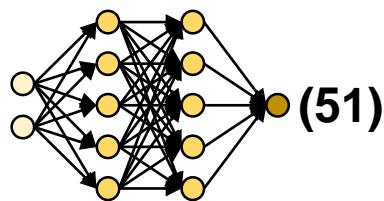
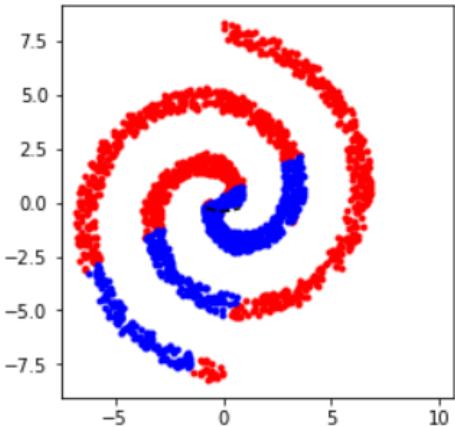
Example n°3 : Clusterization_not_linearly_separated_spirale



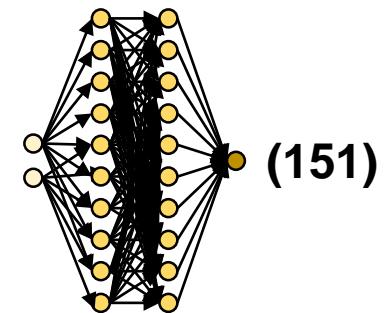
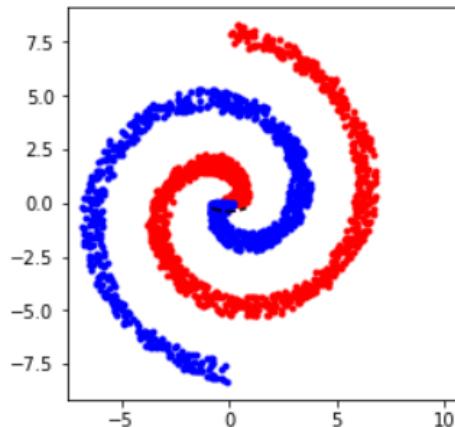
Number of neurons vs. fitting capacity



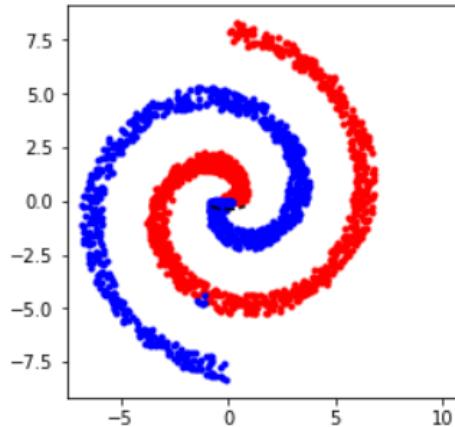
Underfitting



Optimized



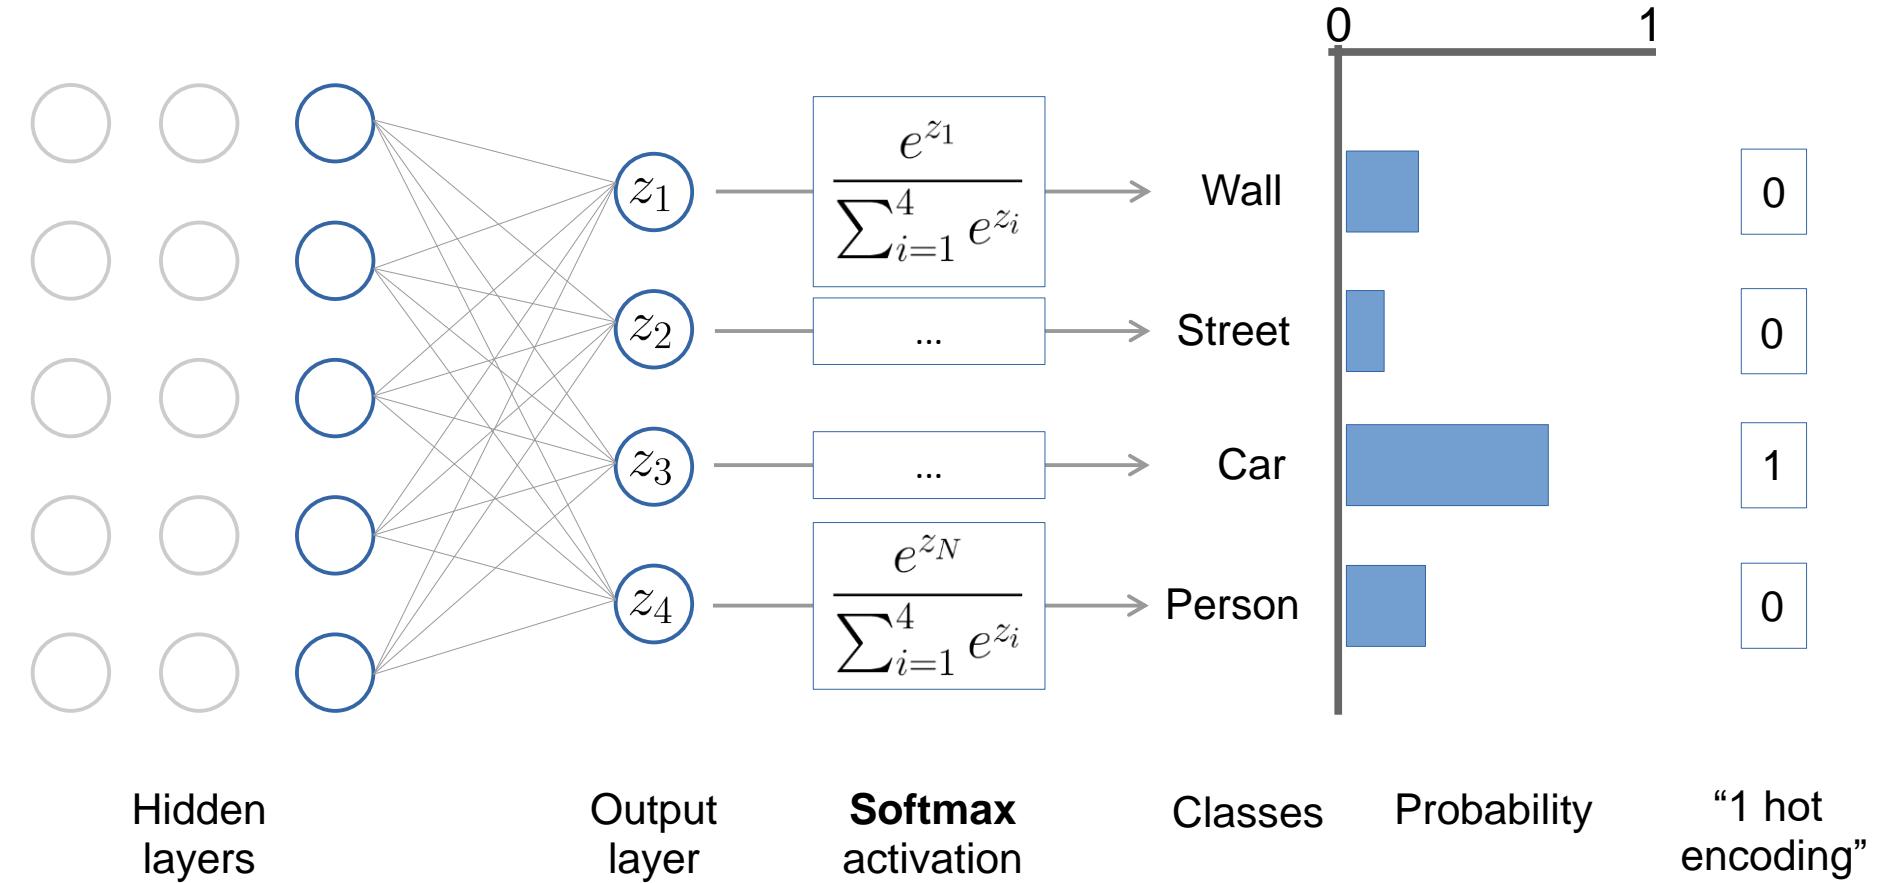
Overfitting



Outline

- I. **Introduction**
- II. **Starting with Deep Learning and Artificial Neuron Network**
 - i. A bit of vocabulary before starting
 - ii. Case 1 : single neuron network
 - iii. Case 2 : how to construct a densely connected network
 - iv. Case 3 : multi-class clusterization**
 - v. Summary

Classifier with multiple classes : softmax activation



Outline

- I. Introduction**
- II. Starting with Deep Learning and Artificial Neuron Network**
 - i. A bit of vocabulary before starting
 - ii. Case 1 : single neuron network
 - iii. Case 2 : how to construct a densely connected network
 - iv. Case 3 : multi-class clusterization
 - v. Summary**

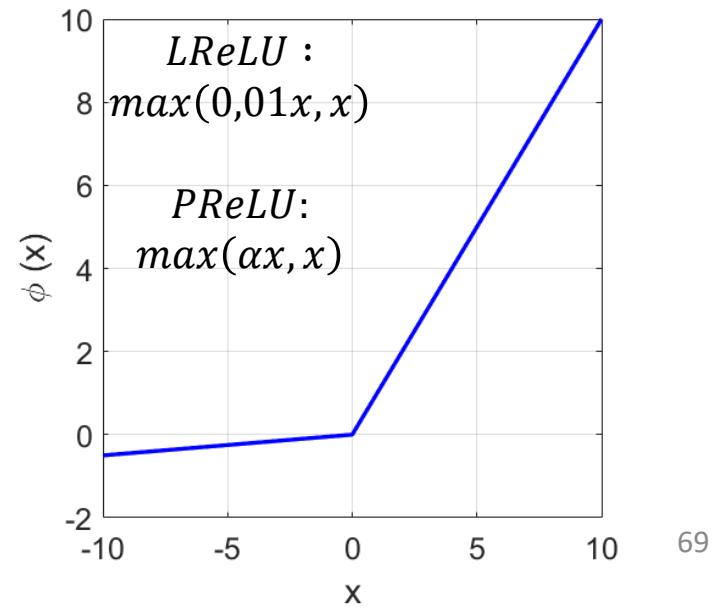
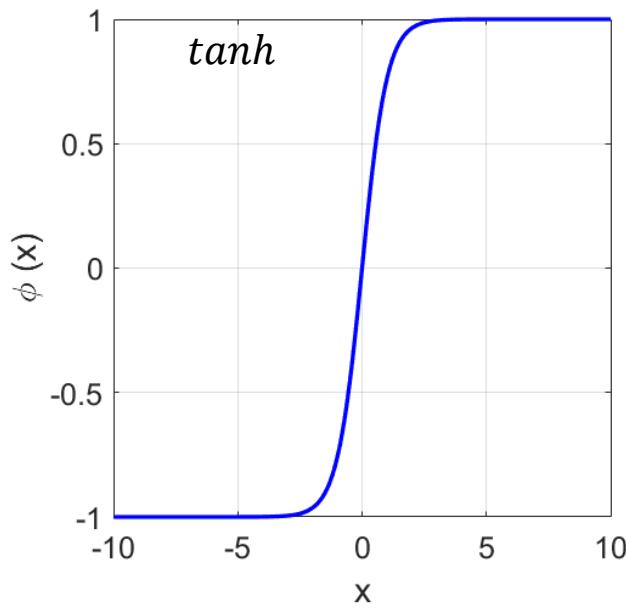
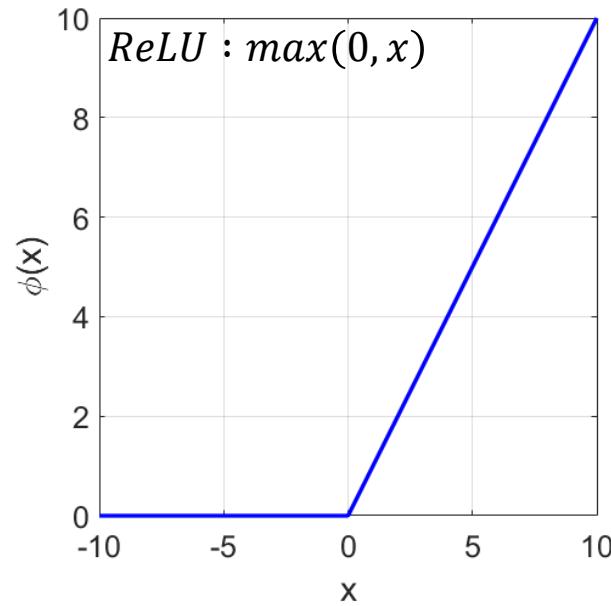
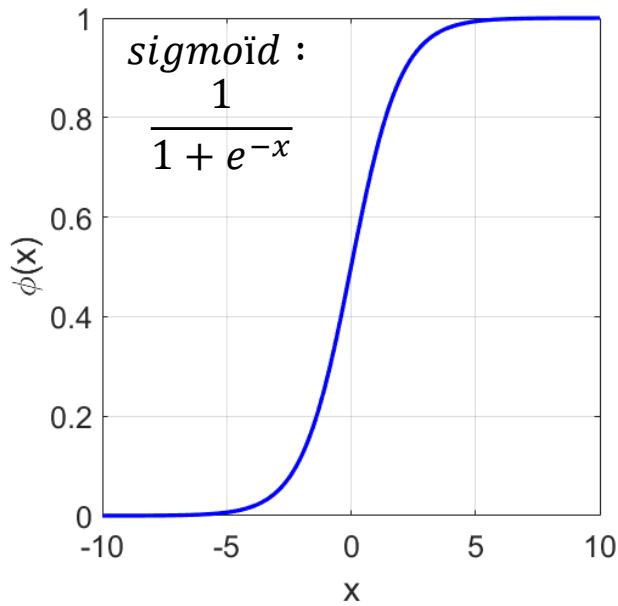
General summary

How to choose the activation & loss function?

Problem type	Last-layer activation	Loss function	Number of neurons in the last layer
Binary classification	'sigmoid'	'binary-crossentropy'	1
Multiclass, single-label classification	'softmax'	'categorical_crossentropy'	As many as the number of classes
Regression to values between 0 and 1	'sigmoid' or 'none'	'mse'	1

As a rule-of-thumbs, use 'relu' everywhere else as activation function

General summary



General summary

Most commonly used optimizers?

Optimize full name	Python call name	Learning-rate
Stochastic Gradient Descent	'sgd'	Fixed by the user
Root Mean Square Propagation	'RMSprop'	Adjusted automatically
Adaptive Moment Optimization	'adam'	Adjusted automatically